

Lab 1 - Introduction.

Google Colab

→ free cloud based Coding environment by google.

Why?

write & run code in python in it.

Supports GPU & TPU

trains dl models faster.

Jupyter Note book

- Open Source tool where we can write code, explain and get op.
- best for data analysis, visualization.
- Runs on your PC not online.

TensorFlow

- dl framework developed by google.
- handles complex neural network building, training & deployment.
- supports both low & high level API.

Keras

- high level API built on top of tensorflow.
- easier & faster than TF.
- Simple syntax.

PyTorch

- dl framework developed by Meta.
- flexible
- debug is made easier.

Numpy

- Handling arrays & mathematical operations.
- Base for most ML libraries.

Pandas

- Data Analysis & Manipulation

Matplotlib

- Plotting graphs & visualizing data.

Scikit-learn

- Traditional ML mode (SVM, DT) & preprocessing.

OpenCV

- Computer Vision task (image & video processing).

Hugging Face Transformers

- NLP tasks & models like BERT, T5PT

Anaconda

- distribution of Python bundled with many libraries, tools & environments used in ML, DL.

Lab 2 : Implement A Classifier Using an Open Source Data Set.

Aim:

To implement a classification model using the SVM algorithm on the IRIS dataset and evaluate its performance using accuracy metrics.

Objectives:

1. To understand + apply the SVM algorithm for classification.
2. Open source dataset (Iris)
3. Preprocess & Split into train & test tests.
4. To train SVM model & make predictions.
5. To evaluate the model's performance using metrics like (accuracy, precision, recall, F1-Score, confusion matrix).

Pseudo Code:

1. Import necessary libraries. (sklearn)
2. Load the Iris data set
3. split the dataset into training & testing.
~~(80% train , 20% test)~~
4. Initialize the SVM classifier.
5. Train the classifier using training data.
6. Predict the o/p. on test data.
7. Evaluate Performance.
8. Display the results.

About the dataset: Iris flower dataset of 150 samples

features \Rightarrow Sepal length, Sepal width, Petal length, Petal width.

Target Label	Flower Species
0	Setosa
1	Versicolor
2	Virginica

Data size: 150 samples (100 training, 50 test)

No. of Samples - 150

No. of features - 4

No. of classes - 3

Samples per class - 50 each.

Why SVM?

\rightarrow best for linearly separable dataset (Setosa)

RBF for Versicolor & Virginica

Observation

Class	Precision	Recall	f1-score
Setosa	1.00	1.00	1.00
Versicolor	1.00	0.93	0.97
Virginica	0.93	1.00	0.97

~~After feature scaling or switching to an RBF Kernel~~

Accuracy: 97.78%

Result: SVM classifier was used on the dataset (RIS) and was implemented successfully.



Scanned with OKEN Scanner

159 lines (159 loc) · 30 KB

[Preview](#) [Code](#) [Blame](#)[Raw](#) [Copy](#) [Download](#)

```
In [14]: import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

# Load dataset
iris = datasets.load_iris()
X, y = iris.data, iris.target

# Standardize features
sc = StandardScaler()
X = sc.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Hyperparameter tuning for RBF SVM
param_grid = {
    'C': [0.1, 1, 10, 100],
    'gamma': [0.01, 0.1, 1, 10],
    'kernel': ['rbf']
}
grid = GridSearchCV(SVC(), param_grid, cv=5)
grid.fit(X_train, y_train)
best_model = grid.best_estimator_

print("Best Parameters:", grid.best_params_)

# Predict on test set
y_pred = best_model.predict(X_test)

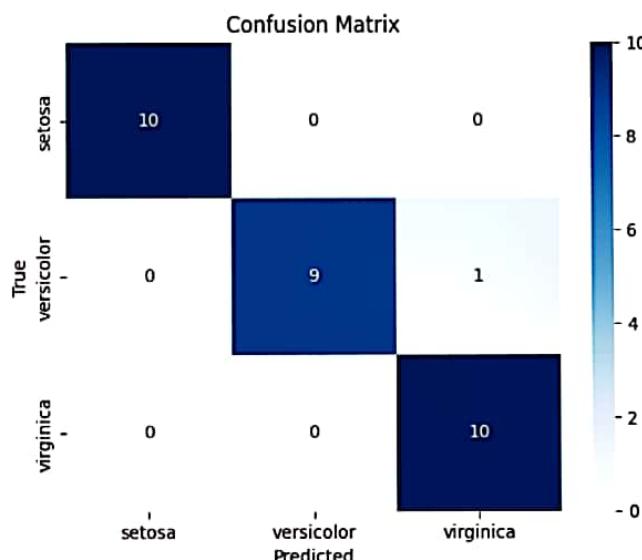
# Evaluation
acc = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {acc*100:.2f}%\n")

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d",
            xticklabels=iris.target_names,
            yticklabels=iris.target_names)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix")
plt.show()

print("\nClassification Report:")
print(classification_report(y_test, y_pred, target_names=iris.target_names))
```

Best Parameters: {'C': 1, 'gamma': 0.1, 'kernel': 'rbf'}

Test Accuracy: 96.67%



Lab 3: Study of the Classifiers with Respect to Statistical Parameters.

Aim:

To study and compare the performance of 2 classification algorithms SVM & Logistic Regression on IRIS dataset Using statistical Evaluation metrics-

Objectives:

1. To implement SVM & Logistic Regression Classifiers
2. To apply both models on the same dataset.
3. To evaluate & compare their performance .
Using Accuracy, Precision, Recall, F1-score.

Data Set Used :

- IRIS
- Sklearn.
- No. of samples : 150 .
No. of features : 4
No. of classes : 3 .

Pseudo Code .

1. Import necessary libraries & load the Iris dataset
2. Scale the data set Using Standard Scaler.
3. Split the dataset into training & testing sets .
4. Initialize both SVM & Logistic Regression Classifiers.
5. Train both classifiers on training data .

- What is statistical Parameters?
- They are numerical values that describe characteristics of a dataset.
 - In ml, it is referred to performance metrics Accuracy, precision, recall, f-score.

Accuracy

- no. of correctly predicted samples out of total.

$$= \frac{\text{No. of Correct Predictions}}{\text{Total No. of Predictions}}$$

Precision

Proportion of true positive out of all positive predictions made by model.

$$= \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Recall / Sensitivity / True (+)ve Rate.

- proportion of true (+)ve out of all actual (+)ves.

$$= \frac{\text{True Positives}}{\text{True Positive} + \text{False Negative.}}$$

F1-Score

harmonic mean of precision and recall.

$$F_1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall.}}$$

6. Predict the test results using both models

7. Evaluate both models using

- Accuracy
- Confusion Matrix.
- Classification Report.

8. Compare the results.

Observation:

Metric	SVM	Logistic Regression.
Accuracy	97.78 %	97.78
Precision	0.98	0.93
Recall	0.98	0.98
F1-Score	0.98	0.98

Both Models performed equally well on the IRIS dataset.

Why Logistic Regression?

- simple & fast.
- training is fast on small datasets.
- works well on linearly separable dataset.

Result:

~~Both~~ Both SVM & Logistic Regression classifiers achieved high performance on IRIS dataset with accuracy 97.78 %.

Thus the experiment was successfully implemented.

main DEEP-LEARNING-TECHNIQUES-LAB-EXP / DLT LAB3.ipynb

Hanishkao01 Add files via upload

233 lines (233 loc) · 119 KB

53eebc0 · 4 minutes ago

Preview Code Blame Raw

```
In [2]:
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix,
import seaborn as sns
import matplotlib.pyplot as plt

# Load dataset
iris = datasets.load_iris()
X, y = iris.data, iris.target

# Standardize features
sc = StandardScaler()
X = sc.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Define classifiers
classifiers = {
    "SVM (RBF)": SVC(kernel='rbf', C=1.0, gamma='scale', random_state=42),
    "Logistic Regression": LogisticRegression(max_iter=200, random_state=42),
    "KNN (k=5)": KNeighborsClassifier(n_neighbors=5),
    "Decision Tree": DecisionTreeClassifier(random_state=42)
}

# Compare classifiers
for name, model in classifiers.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='macro')
    rec = recall_score(y_test, y_pred, average='macro')
    f1 = f1_score(y_test, y_pred, average='macro')

    print(f"\n{name} Statistics:")
    print(f"Accuracy : {acc*100:.2f}%")
    print(f"Precision: {prec*100:.2f}%")
    print(f"Recall   : {rec*100:.2f}%")
    print(f"F1-Score : {f1*100:.2f}%")

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d",
            xticklabels=iris.target_names,
            yticklabels=iris.target_names)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title(f"{name} - Confusion Matrix")
plt.show()

# Detailed Classification Report
print("Classification Report:\n", classification_report(y_test, y_pred, target_names=iris.target_ni)
```

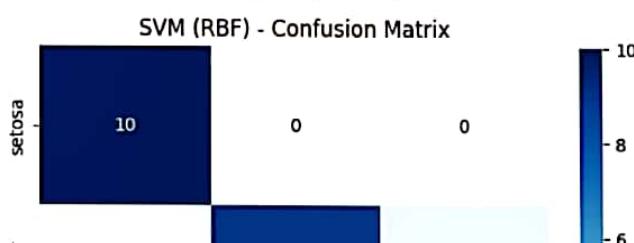
SVM (RBF) Statistics:

Accuracy : 96.67%

Precision: 96.97%

Recall : 96.67%

F1-Score : 96.66%



Ex:4 Build a Simple feed forward neural network to recognise handwritten characters.

Aim:

To design, implement and evaluate a simple feedforward neural network for recognising handwritten characters.

Objectives:

1. Understand the architecture of a feed forward neural network.
2. To apply supervised learning for image classification.
3. To preprocess the MNIST dataset for neural network training.
4. To evaluate the model's accuracy & analyse the performance.

Algorithm:

1. Import libraries - pytorch, Keras, numpy.
2. Load dataset MNIST
3. Preprocess the data. normalise, flatten, convert (0-1) 28×28 labels.
4. Define model.
Input layer: 784 neurons
Hidden layer: 528 neurons
O/P layer: 10 neurons
5. Compile model.
6. Train model.
7. Evaluate model.
8. Predict & visualise.

PseudoCode :

START

Import required libraries.

Load MNIST dataset

Normalise pixel values between 0 & 1

Flatten 28x28 images into vectors.

One-hot encode the labels.

Define a sequential Neural network.

Input layer : size 784.

Hidden layer : 128 neurons,

Compile model with Adam optimizer

Softmax activation
Categorical cross entropy loss, accuracy metric.

Train model using training dataset for define epochs.

Evaluate model on test dataset.

Display accuracy + sample predictions

END.

Observations:

- achieved 97.96% accuracy on MNIST dataset with 5 epochs of training.
- Increasing hidden units can improve performance.
- Model might struggle with similar characters.
- Error with centered, grayscale, similar stroke thickness.
- Invert colours.

Result:

The mode was successfully executed.

main · DEEP-LEARNING-TECHNIQUES AND EX / DEEPLearnPyTorch

Hanishka01 Add files via upload 5407206 · 40 minutes ago

414 lines (414 loc) · 52.9 KB

Preview Code Blame Raw

```
In [1]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms

# Load dataset
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,), (0.5,))])

trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=64, shuffle=True)

testset = torchvision.datasets.MNIST(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=64, shuffle=False)

# Define Feedforward Neural Network
class FeedForwardNN(nn.Module):
    def __init__(self):
        super(FeedForwardNN, self).__init__()
        self.fc1 = nn.Linear(28*28, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 10)
        self.relu = nn.ReLU()
    def forward(self, x):
        x = x.view(-1, 28*28)
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.fc3(x)
        return x

model = FeedForwardNN()

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training
epochs = 5
for epoch in range(epochs):
    running_loss = 0.0
    for images, labels in trainloader:
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
    print(f"Epoch {epoch+1}, Loss: {running_loss/len(trainloader):.4f}")

# Testing
correct, total = 0, 0
with torch.no_grad():
    for images, labels in testloader:
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f"Accuracy on test set: {100 * correct / total:.2f}%")
```



Epoch 1, Loss: 0.4191
Epoch 2, Loss: 0.1967
Epoch 3, Loss: 0.1437
Epoch 4, Loss: 0.1151
Epoch 5, Loss: 0.1007
Accuracy on test set: 96.16%

```
In [5]: from PIL import Image
import torchvision.transforms as transforms
import torch
from google.colab import files
import matplotlib.pyplot as plt

# Load trained model
model.eval()

# Preprocessing transform (same as training)
transform = transforms.Compose([
    transforms.Resize((28,28)),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))])
```



D Pumping Lemma

If a language L is Regular, pumping length P ,

Any string $w \in L$ of length $|w| \geq P$

Consider $w = xyz$

where i) $|xy| \leq P$

ii) $|y| \geq 1$

iii) For All $i \geq 0$, $yyz \in L$

i) $L_1 = \{w \in \{a, b\}^* \mid \text{no. of } a's \text{ in } w = \text{no. of } b's\}$

Assume L_1 is regular

$L_1 = \{ab, aabb, aaabbb, \dots\}$

Let $w = aabb$ and $P = 4$

$|w|=4$, $\therefore |w| \geq P$

Consider $x=a$, $y=ab$, $z=4$

$$|xy| \leq P$$

$$2 \leq 4$$

$$|y| \geq 1$$

$$z \geq 1$$

$xy^iz \in L_1$, for all $i \geq 0$

$$i=0 \rightarrow xz \rightarrow ab$$

$$i=1 \rightarrow xy^i z \rightarrow aabb$$

$$i=2 \rightarrow xy^{i-1} z \rightarrow aabbab \notin L_1$$

Contradiction

L_1 is not Regular.

ii) $L_2 = \{w \text{ in } \{a,b\}^* \mid w \text{ is a Palindrome}\}$

Assume L_2 is Regular

$L_2 = \{aba, aabaa, aaabaaa, \dots\}$

Let $w = aabaa$ & $P = 5$: $|w| \geq P$

Consider $x = aa$, $y = b$, $z = aa$

$$|x| \leq P$$

$$4 \leq 5$$

$$1 \geq 1$$

$$|y| \geq 1$$

$$1 \geq 1$$

$$\forall i \geq 0 \rightarrow x^i z \in L_2$$

$$i=0, xz = aaaa \notin L_2$$

$$i=1, xyz = aabaa$$

L_2 is not Regular

iii) $L_3 = \{a^k \mid k \text{ is Prime Number}\}$

Assume L_3 is Regular

$L_3 = \{aa, aaa, aaaga, \dots\}$

Let $w = aaa$ & $P = 3$

$$|w| = 3 \therefore |w| \geq P$$

Consider $x = a$, $y = a$, $z = a$

$$|xy| \leq P$$

$$2 \leq 3$$

$$|y| \geq 1$$

$$1 \geq 1$$

$$\forall i \geq 0, x^i y^i z \in L_3$$

$$i=0, yz = aa$$

$$i=1, yyz = aaa$$

$$i=2, y^2 z = aaaa \notin L_3$$

$\Rightarrow L_3$ is not Regular

iv) $L_4 = \{a^k \mid k \text{ is a perfect square}\}$

Assume L_4 is Regular

$$L_4 = \{a, aa, aaa, \dots\}$$

Let $w = aaaa$ and $P=4$

$$|w|=4 \therefore |w| \geq P$$

Consider $x=a$, $y=aa$, $z=a^4$

$$|xy| \leq P$$

$$3 \leq 4$$

$$|y| \geq 1$$

$$2 \geq 1$$

$$\forall i \geq 0, xy^i z \in L_4$$

$$i=0, yz = aaa \notin L_4$$

L_4 is not Regular

v) $L_5 = \{a^k \mid k \text{ is a factorial number}\}$

Assume L_5 is Regular

$$L_5 = \{a, aa, aaaaa, a^2, \dots\}$$

Let $w = aaaaaaaaa$ and $P=6$

$$|w|=6 \therefore |w| \geq P$$

Consider $x=a^3$, $y=aa$, $z=a^3$

$$|xy| \leq P$$

$$4 \leq 6$$

$$|y| \geq 1$$

$$|z| \geq 1$$

$$\forall i \geq 0, xy^i z \in L_5$$

$$i=0, yz = aaaa \notin L_5, \text{ Contradiction}$$

$$i=1, xyz = aaa aaa \in L_5$$

L_5 is not Regular.

- 2) If L and R are two Regular languages
- L concatenated with $[L \cap R]$
- $L \cdot (L \cap R)$
 given L and R are Regular
 By Closure Properties
 Intersection of Regular languages is Regular
 $L \cap R$ is Regular
- Concatenation of Regular languages is Regular
 $L \cdot (L \cap R)$ is Regular
 $\therefore L \cdot (L \cap R)$ is Regular
- Complement of R union (Reversal of L)
 L & R are Regular
 By Closure Properties.
- Complement of Regular lang is Regular (R^c is Regular)
 - Reversal of Regular language is Regular (L^R is Regular)
 - Union of Regular language is Regular ($R \cup L$ is Regular)
- $\therefore (R^c) \cup (L^R)$ is Regular.
- (Kleene star of L) set difference (Kleene star of R) $(L^*) \setminus (R^*)$
 given L & R are Regular.
 - By Closure Properties
 - Kleene star of Regular language is Regular (L^* is Regular)
 - Complement of Regular language is Regular (R^c is Regular)
 - Intersection of Regular lang is Regular ($L \cap R$ is Regular)

$$(L^*) \setminus (R^*) = L^* \cap (R^*)^c$$

$\therefore (L^*) \setminus (R^*)$ is Regular.

Lab5: Study Of Activation function and their Role.

Aim:

To study the different activation functions in dl, understand their role & analyse their effects on training & performance of NN

Objectives:

- Understand the purpose of activation func in ml.
- implement various activation functions such as Sigmoid, Tanh, ReLU & Leaky ReLU.
- Analyse and Compare the performance of a simple neural network using different Activation funct.
- Observe the effects of activation functions on learning dynamics.
- draw conclusion on best activation function.

Common Activation functions:

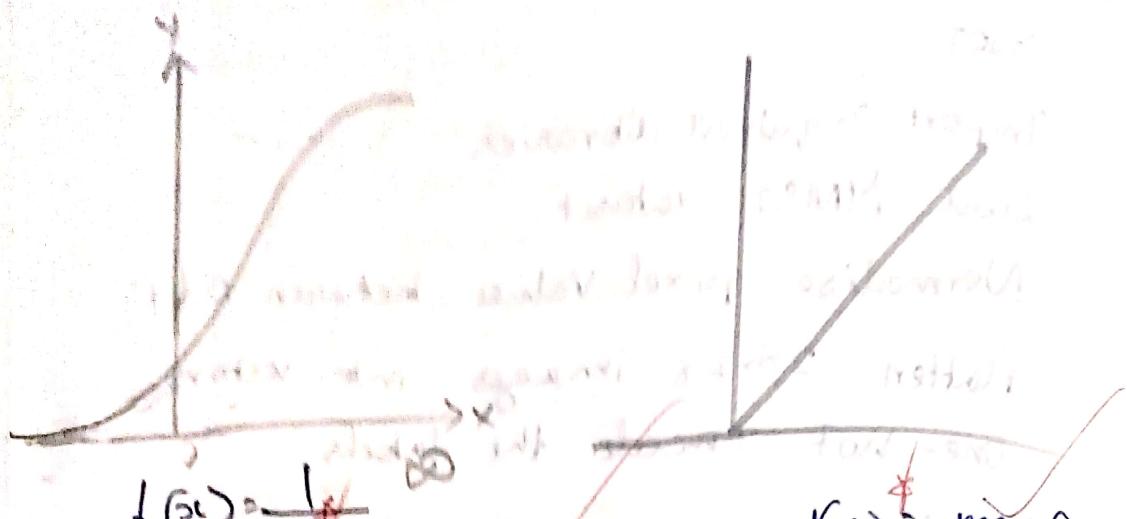
Sigmoid	- Squashes input values to $[0, 1]$
Tanh	- Squashes input Values $[-1, 1]$
ReLU	- Outputs 0 for negative inputs & linear for +ve inputs.
Leaky ReLU	- Allows a small, non zero gradients when ip is -ve.

Algorithm:

1. Install required libraries (numpy, matplotlib)
2. Implement Python code to define and compute activation functions.
3. Generate input data using Numpy.
4. Apply activation function and Visualise their

sigmoid

ReLU



$$f(x) = \frac{1}{1 + e^{-x}}$$

activation function for sigmoid

range $\rightarrow (0, 1)$

$$f(x) = \max(0, x)$$

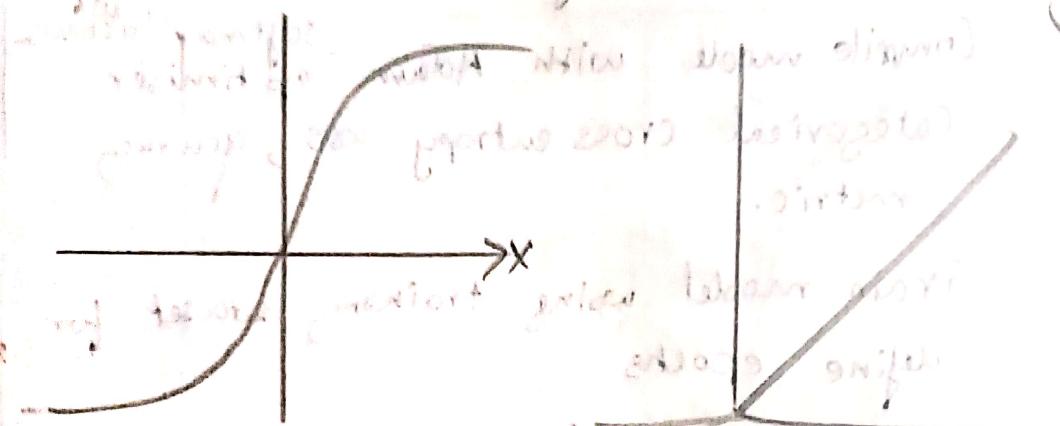
dead neuron

tanh

V.G.

$(0, \infty)$

Leaky ReLU



$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$(-1, 1)$

$$f(x) = 0.01x \quad x > 0$$

$0.1x$

$$V.G. \text{ but slightly better: } f(x) = \max(0.1x, x)$$

$(-\infty, \infty)$

outputs using Matplotlib.
3. Analyse the behaviour of each function.

Pseudo Code:

1. Initialize

- Training dataset (x, y)

Initialize weights ($w_1, w_2 \dots$) & biases ($b_1, b_2 \dots$) randomly

Choose activation function : ReLU / Sigmoid / Tanh

Set learning rate & number of epochs.

2. Training loop-

for each epoch .

for each training sample (x, y):

$$z_1 = w_1 * x + b_1$$

$$a_1 = \text{Activation}(z_1)$$

$$z_2 = w_2 * a_1 + b_2$$

$$a_2 = \text{Activation}(z_2)$$

$$z_3 = w_3 * a_2 + b_3$$

$$y_{\text{pred}} = \text{Softmax}(z_3)$$

$$\text{loss} = \text{CrossEntropy}(y, y_{\text{pred}})$$

Compute gradients of loss w.r.t w_1, w_2, w_3 & b_1, b_2, b_3 .

Update weights & biases using gradient descent .

END FOR

Print current epoch & loss .

3. Testing

input test data

Perform forward pass to get prediction

Compare predicted class vs actual label

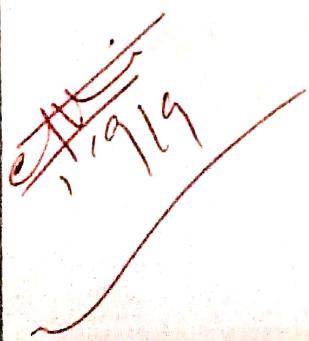
Compute accuracy

END



Observations:

1. Sigmoid: Outputs are compressed b/w 0 & 1,
suitable for probability based O/P
but prone to vanishing gradients.
2. ReLU: O/P 0 for c-sve I/P, promoting
sparsity but may cause "dying ReLU",
where neurons O/P 0 permanently.
3. Tanh: O/P range from -1 to 1,
useful for centered data, but also
suffers from vanishing gradient for
large inputs.
4. Leaky ReLU: Allows small c-sve o/p,
mitigating the "dying ReLU" problem.



Result:

The code was successfully executed
and verified.

```

In [2]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,),

trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True, tra
testset = torchvision.datasets.MNIST(root='./data', train=False, download=True, tra

trainloader = torch.utils.data.DataLoader(trainset, batch_size=64, shuffle=True)
testloader = torch.utils.data.DataLoader(testset, batch_size=1000, shuffle=False)

class SimpleNN(nn.Module):
    def __init__(self, activation_fn):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(28*28, 256)
        self.fc2 = nn.Linear(256, 128)
        self.fc3 = nn.Linear(128, 10)
        self.activation = activation_fn

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = self.activation(self.fc1(x))
        x = self.activation(self.fc2(x))
        return self.fc3(x)

def train_and_eval(activation_fn, epochs=5):
    model = SimpleNN(activation_fn)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.001)

    train_losses = []
    test_accuracies = []

    for epoch in range(epochs):
        model.train()
        running_loss = 0.0
        for images, labels in trainloader:
            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

        avg_loss = running_loss / len(trainloader)
        train_losses.append(avg_loss)

        # Evaluate on test set
        model.eval()
        correct, total = 0, 0
        with torch.no_grad():
            for images, labels in testloader:
                outputs = model(images)
                _, predicted = torch.max(outputs, 1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()
        acc = correct / total
        test_accuracies.append(acc)

        print(f"Epoch [{epoch+1}/{epochs}], Loss: {avg_loss:.4f}, Test Acc: {acc*10

    return train_losses, test_accuracies

activations = {
    "Sigmoid": nn.Sigmoid(),
    "Tanh": nn.Tanh(),
    "ReLU": nn.ReLU(),
    "LeakyReLU": nn.LeakyReLU(0.1)
}

results = {}
for name, act in activations.items():
    print(f"\nTraining with {name} activation...")
    train_losses, test_accuracies = train_and_eval(act, epochs=5)
    results[name] = (train_losses, test_accuracies)

plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
for name, (losses, accs) in results.items():
    plt.plot(losses, label=name)
plt.title("Training Loss vs Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.subplot(1,2,2)
for name, (losses, accs) in results.items():

```



Lab 6: Implement Gradient Descent And Backpropagation in Deep Neural Network.

Aim:

To implement and understand the working of gradient descent and backpropagation in training a deep neural network for supervised learning tasks.

Objectives:

- how weights & biases are updated using G.D.
- forward pass & backward pass in DNN.
- minimize loss function using optimisation.
- learning rate & iterations affect convergence.
- trained network's predictions compared to expected o/p.

Algorithm:

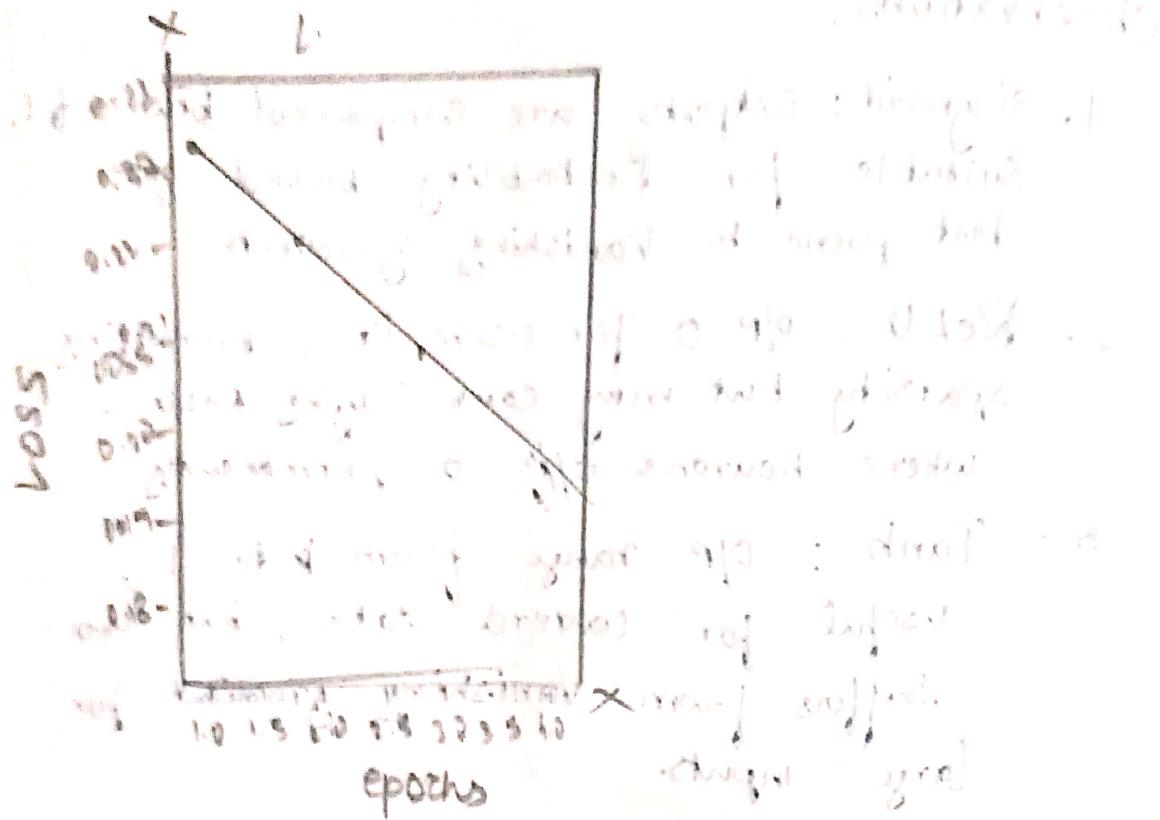
- 1) Initialize network architecture (no. of i/p, hidden, o/p nodes)
- 2) Randomly initialize weights & biases.
- 3) Forward propagation.
 - input \rightarrow hidden using weighted sum.
 - Apply activation func.
 - Compute o/p layer values.

Loss Calculation.

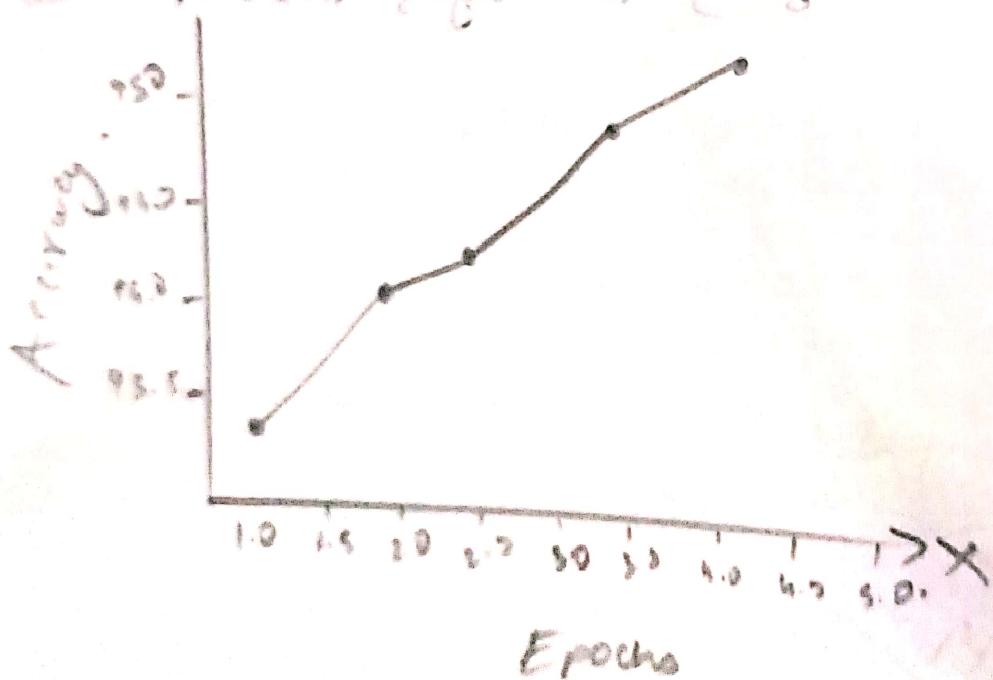
- Compute error using a loss function (MSE/CE)

Backward Propagation.

- Compute gradient of loss w.r.t. weights & biases.
- Update weights & biases using G.D rule.



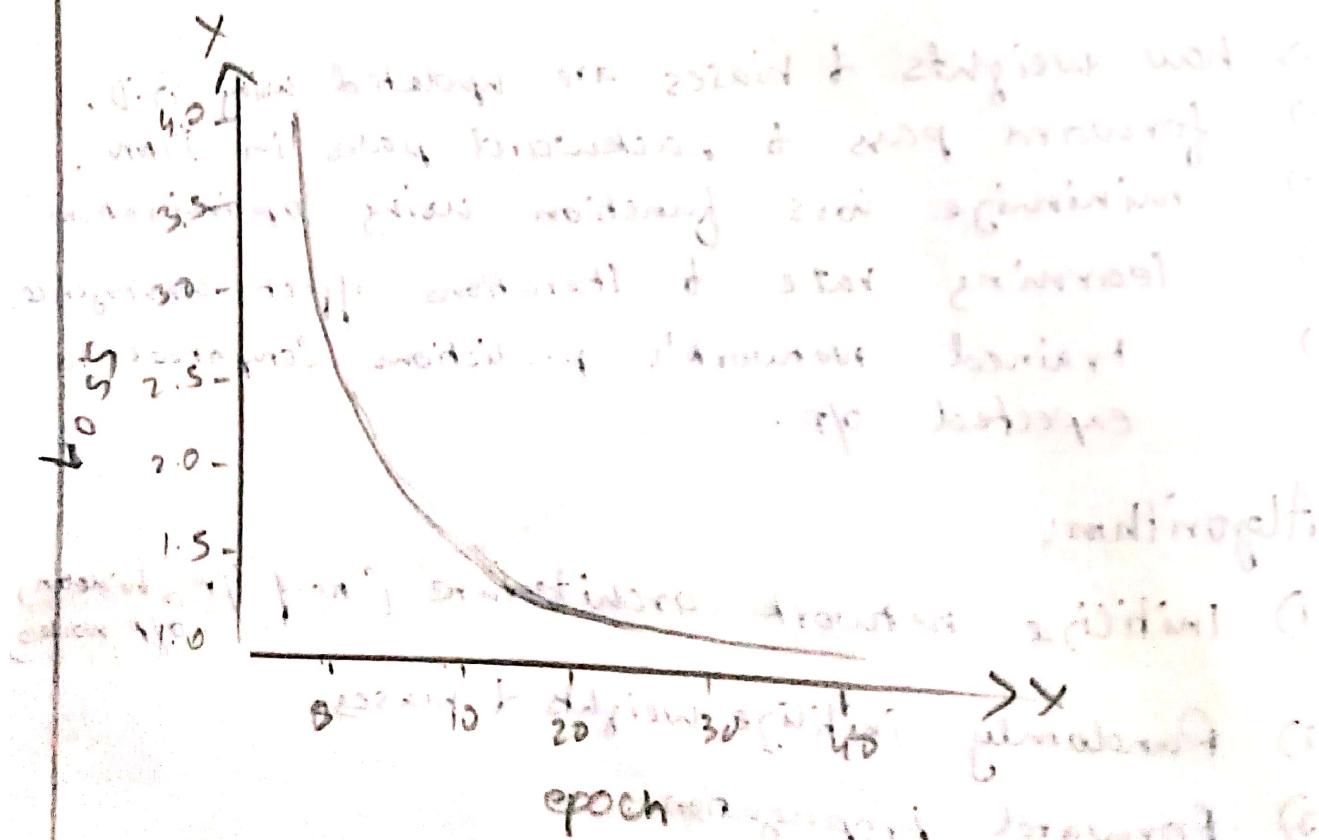
The loss has fallen to 0.02 after 10 epochs of training.



The accuracy has increased to 94.8% after 10 epochs of training.

Epoch 1 : LOSS : 0.2449 Acc : 93.14%
 Epoch 2 : LOSS : 0.2246 Acc : 94.00%
 Epoch 3 : LOSS : 0.2059 Acc : 94.25%
 Epoch 4 : LOSS : 0.1893 Acc : 94.83%
 Epoch 5 : LOSS : 0.1750 Acc : 95.23%

loss curve



$$W = W - n \frac{\partial L}{\partial W}$$

n → learning rate.

$$b = b - n \frac{\partial L}{\partial b}$$

- 4) Repeat until convergence
- 5) O/P trained weights , loss trend

PseudoCode:

Initialize weights W , biases b , learning-rate, epochs
for each epoch :

$$\text{output} = \text{forward-pass}(x, -W, b)$$

$$\text{loss} \leftarrow \text{complete-loss}(y_{\text{true}}, \text{output})$$

$$\text{gradients} = \text{backward-pass}(\text{output}, y_{\text{true}}, W, b)$$

$$W, b = \text{update-parameters}(W, b, \text{gradients}, \text{learning rate})$$

Observation:

- loss decreases as epochs increases
- Small L.R → slow convergence but stable.
- Large L.R → faster convergence but risk of divergence.
- Network accuracy improves with proper tuning of L.R & hidden layers.

Result:

The code was successfully executed & verified.

```

In [1]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
from tqdm import trange

def sigmoid(x): return 1.0 / (1.0 + np.exp(-x))
def dsigmoid(x):
    s = sigmoid(x)
    return s * (1 - s)

def tanh(x): return np.tanh(x)
def dtanh(x):
    return 1.0 - np.tanh(x)**2

def relu(x): return np.maximum(0, x)
def drelu(x): return (x > 0).astype(float)

def leaky_relu(x, a=0.01):
    return np.where(x > 0, x, a * x)
def dleaky_relu(x, a=0.01):
    dx = np.ones_like(x)
    dx[x <= 0] = a
    return dx

def softmax(z):
    z = z - np.max(z, axis=1, keepdims=True)
    expz = np.exp(z)
    return expz / np.sum(expz, axis=1, keepdims=True)

def cross_entropy_loss(probs, y_true_onehot):
    m = y_true_onehot.shape[0]
    p = np.clip(probs, 1e-12, 1.0)
    return -np.sum(y_true_onehot * np.log(p)) / m

class MLP:
    def __init__(self, input_dim, hidden_dims, output_dim, activation='relu', lr=0.1, weight_scale=0.01):
        np.random.seed(seed)
        self.layers = [input_dim] + hidden_dims + [output_dim]
        self.L = len(self.layers) - 1
        self.lr = lr

        self.W = [None] * self.L
        self.b = [None] * self.L
        for i in range(self.L):
            self.W[i] = np.random.randn(self.layers[i], self.layers[i+1]) * weight_scale
            self.b[i] = np.zeros((1, self.layers[i+1]))

        if activation == 'sigmoid':
            self.act = sigmoid; self.dact = dsigmoid
        elif activation == 'tanh':
            self.act = tanh; self.dact = dtanh
        elif activation == 'relu':
            self.act = relu; self.dact = drelu
        elif activation == 'leaky':
            self.act = lambda x: leaky_relu(x, a=0.01)
            self.dact = lambda x: dleaky_relu(x, a=0.01)
        else:
            raise ValueError("activation must be: sigmoid, tanh, relu, leaky")

    def forward(self, X):
        """Forward pass. Stores pre-activations (Z) and activations (A)."""
        A = X
        self.As = [A]
        self.Zs = []
        for i in range(self.L - 1):
            Z = A.dot(self.W[i]) + self.b[i]
            A = self.act(Z)
            self.Zs.append(Z)
            self.As.append(A)

        Z = A.dot(self.W[-1]) + self.b[-1]
        self.Zs.append(Z)
        self.As.append(Z)
        probs = softmax(Z)
        return probs

    def backward(self, probs, y_onehot):
        """Backpropagation. Compute gradients for all W and b."""
        m = y_onehot.shape[0]
        grads_W = [None] * self.L
        grads_b = [None] * self.L

        dZ = (probs - y_onehot) / m
        grads_W[-1] = self.As[-2].T.dot(dZ)
        grads_b[-1] = np.sum(dZ, axis=0, keepdims=True)

        dA_prev = dZ.dot(self.W[-1].T)
        for l in range(self.L - 2, -1, -1):
            Z = self.Zs[l]
            dZ = dA_prev * self.dact(Z)
            grads_W[l] = self.As[l].T.dot(dZ)
            grads_b[l] = np.sum(dZ, axis=0, keepdims=True)
            if l > 0:
                dA_prev = dZ.dot(self.W[l].T)

        return grads_W, grads_b

    def step(self, grads_W, grads_b):
        """SGD parameter update (vanilla)."""
        for i in range(self.L):
            self.W[i] -= self.lr * grads_W[i]
            self.b[i] -= self.lr * grads_b[i]

    def predict(self, X):
        probs = self.forward(X)

```



```

def predict(self, X):
    probs = self.forward(X)
    return np.argmax(probs, axis=1)

(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype(np.float32) / 255.0
x_test = x_test.astype(np.float32) / 255.0
x_train = x_train.reshape(-1, 28*28)
x_test = x_test.reshape(-1, 28*28)

y_train_oh = to_categorical(y_train, 10)
y_test_oh = to_categorical(y_test, 10)

def train_experiment(activation='relu',
                     hidden_dims=[128, 64],
                     lr=0.1,
                     epochs=10,
                     batch_size=128,
                     weight_scale=0.05,
                     seed=123):
    model = MLP(input_dim=28*28, hidden_dims=hidden_dims, output_dim=10,
                activation=activation, lr=lr, weight_scale=weight_scale, seed=seed)
    history = {'train_loss': [], 'train_acc': [], 'test_loss': [], 'test_acc': []}
    n = x_train.shape[0]
    steps_per_epoch = n // batch_size

    for epoch in range(epochs):
        perm = np.random.permutation(n)
        x_shuf = x_train[perm]
        y_shuf = y_train_oh[perm]
        epoch_loss = 0.0

        for i in range(steps_per_epoch):
            start = i * batch_size
            end = start + batch_size
            Xb = x_shuf[start:end]
            yb = y_shuf[start:end]

            probs = model.forward(Xb)
            loss = cross_entropy_loss(probs, yb)
            epoch_loss += loss

            grads_W, grads_b = model.backward(probs, yb)

            model.step(grads_W, grads_b)

        avg_loss = epoch_loss / steps_per_epoch
        y_pred_train = model.predict(x_train[:2000])
        train_acc = (y_pred_train == y_train[:2000]).mean()

        probs_test = model.forward(x_test)
        test_loss = cross_entropy_loss(probs_test, y_test_oh)
        test_pred = np.argmax(probs_test, axis=1)
        test_acc = (test_pred == y_test).mean()

        history['train_loss'].append(avg_loss)
        history['train_acc'].append(train_acc)
        history['test_loss'].append(test_loss)
        history['test_acc'].append(test_acc)

    print(f"Epoch {epoch+1}/{epochs} | train_loss={avg_loss:.4f} train_acc={train_acc*100:.2f}% | "
          "test_loss={test_loss:.4f} test_acc={test_acc*100:.2f}%")

    return model, history

activations = ['sigmoid', 'tanh', 'relu', 'leaky']
results = {}

for act in activations:
    print("\n" + "*50)
    print(f"Training with activation: {act}")
    model, hist = train_experiment(activation=act, hidden_dims=[256, 128], lr=0.1, epochs=8, batch_size=128)
    results[act] = hist

for act, hist in results.items():
    plt.figure(figsize=(10,4))
    plt.subplot(1,2,1)
    plt.plot(hist['train_loss'], marker='o')
    plt.title(f"{act} - Train Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")

    plt.subplot(1,2,2)
    plt.plot(hist['test_acc'], marker='o')
    plt.title(f"{act} - Test Accuracy")
    plt.xlabel("Epochs")
    plt.ylabel("Accuracy")

    plt.suptitle(f"Activation: {act}")
    plt.show()

```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 
0s 0us/step 



Scanned with OKEN Scanner

Lab 7: BUILD A CNN MODEL TO CLASSIFY CATS AND DOGS DATASET.

Aim:

To build and train a CNN model using transfer learning to classify images of cats & dogs with high accuracy.

Objective

- Understand and apply CNN for image classification
- Utilize transfer learning with pretrained model
- Implement data preprocessing and augmentation
- Evaluate model performance on unseen test data
- Test the trained model on new images.

Algorithm

1. Download & extract the cats & dogs dataset
2. Prepare data generators
 - Apply data augmentation & preproccesing
 - Split training data into training & validation sets.
3. Load pretrained MobileNetV2 model without top layers.
4. Add custom classification head:
 - Global Avg Pool.
 - Dropout layer for regularisation.
 - Dense opp layer with Sigmoid Activation for binary Classification.

Layer Type	Output Shape	Param
Conv2d-8 (Conv2D)	(None, 16, 16, 32)	896
Max pooling2d-8 (MaxPooling2D)	(None, 8, 8, 32)	0

Max pooling2d-8 (None, 8, 8, 32) Output shape: 0
 (MaxPooling2D) Stage: 2 Epoch: 11 Loss: 0.5812 Val loss: 0.6211

Conv2d-9 (Conv2D) (None, 7, 7, 64) 18,496
 (MaxPooling2D) Stage: 3 Epoch: 12 Loss: 0.5812 Val loss: 0.6211

Max pooling2d-9 (None, 4, 4, 64) Stage: 3 Epoch: 12 Loss: 0.5812 Val loss: 0.6211

Conv2d-10 (Conv2D) (None, 3, 3, 128) 73,856
 (MaxPooling2D) Stage: 4 Epoch: 13 Loss: 0.5812 Val loss: 0.6211

Max Pooling2d-10 (None, 1, 1, 128) 0
 (MaxPooling2D) Stage: 4 Epoch: 13 Loss: 0.5812 Val loss: 0.6211

Epoch 1: acc 0.5092 loss: 0.7737 val.acc 0.411
 val loss: 0.6211

Epoch 2: acc: 0.5188 loss: 0.6858 val.acc: 0.6858 val loss: 0.6211

Epoch 3: acc: 0.4794 loss: 0.6928 val.acc: 0.4688 val loss: 0.6211

Epoch 4: acc: 0.4062 loss: 0.6933 val.acc: 0.500 val loss: 0.6210

Epoch 5: acc: 0.5055 loss: 0.6912 val.acc: 0.6250 val loss: 0.6210

Epoch 6: acc: 0.6250 loss: 0.6724 val.acc: 0.6042 val loss: 0.6210

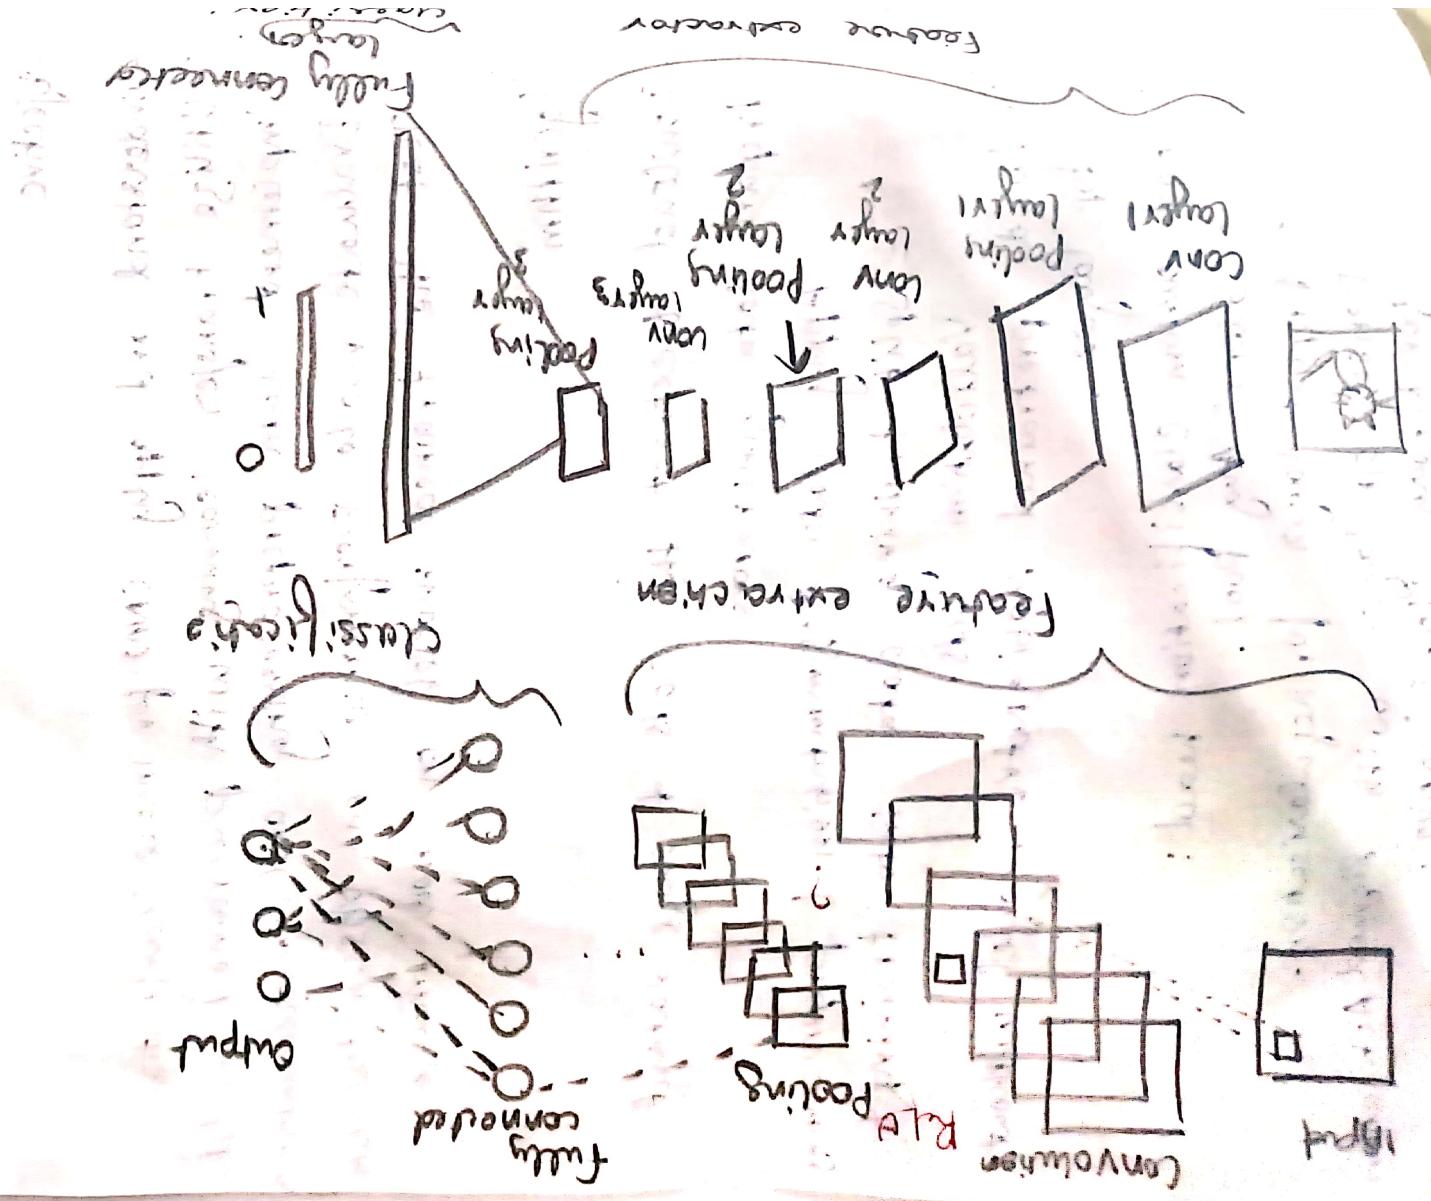
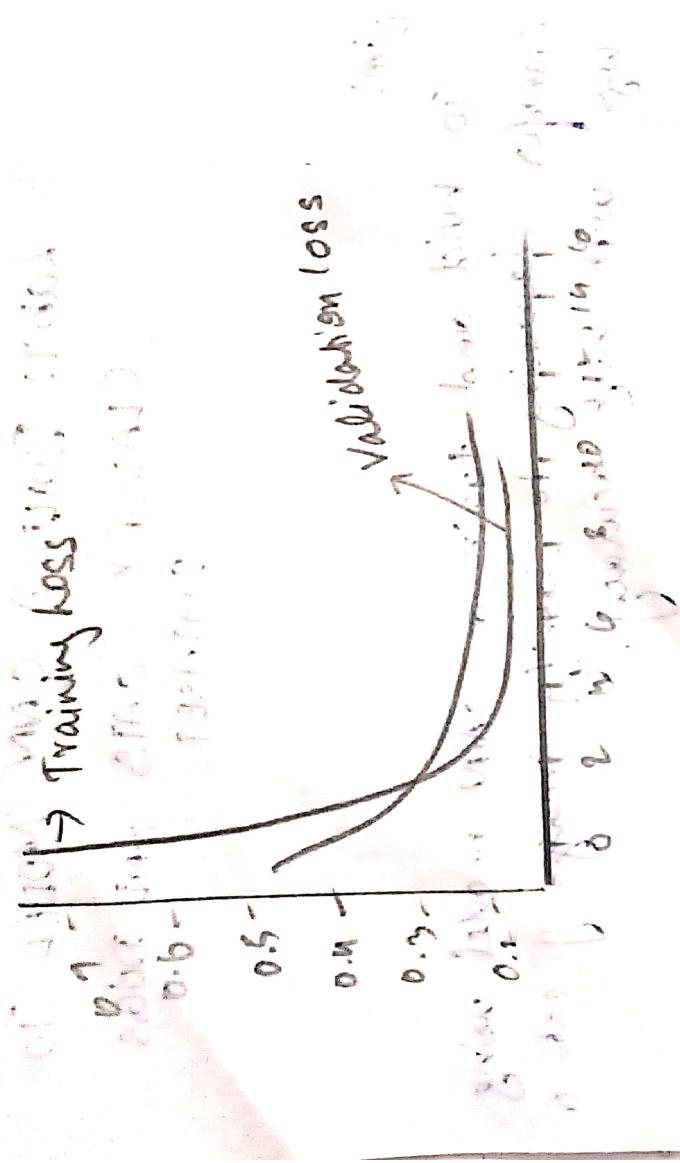
Epoch 7: acc: 0.5952 loss: 0.6875 val.acc: 0.6042 val loss: 0.6210

Epoch 8: acc: 0.6250 loss: 0.6633 val.acc: 0.6146 val loss: 0.6210

Epoch 9: acc: 0.6452 loss: 0.6038 val.acc: 0.6346 val loss: 0.6210

Epoch 10: acc: 0.7500 loss: 0.5812 val.acc: 0.6562 val loss: 0.6210

1



5. Compile a model with adam optimiser & binary cross entropy loss.
6. Train the model on training data with validation.
7. Unfreeze last few layers of base model for fine tuning -
8. Evaluate the model on test data.
9. Save the trained model.
10. Load saved model & preprocess new jpg images.

Observation :

- Pretrained model helped increase accuracy.
- Data augmentation helped reduce overfitting.
- Final test accuracy reached upto 90-95%.
- Model could relatively predict on new unseen images with high confidence.

Result ..

The code was verified successfully and executed

8. Build a LSTM

Aims:

To build a Long Term Memory model that Predicts future Stock Prices based on historical data.

Objective:

1. Develop a deep learning model using LSTM
2. Understand how LSTM networks can capture long term dependencies in time series data.
3. Compare the predicted stock prices with actual prices to evaluate the model's accuracy.
4. Demonstrate that LSTM perform better than simple RNN in regression based sequence prediction Tasks.

Pseudocode:

1. Start
2. Import necessary libraries
numpy, pandas, matplotlib, tensorflow
sklearn, yfinance.
~~import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.preprocessing import timeseriesdataset
from tensorflow.keras.preprocessing import timeseriesdataset~~
3. Load dataset.
 - download Google stock data using yfinance
select the close price column.
4. Preprocess the data
 - Normalise data between 0 & 1 using MinMaxScaler
 - Split dataset into training(80%) & testing(20%)

Layer: Sequential Model

Layer (type)	output shape	Param #
LSTM	(None, 60, 50)	10,400
dropout (Dropout)	(None, 60, 50)	0
lstm_1 (LSTM)	(None, 50)	20,200
dropout_1 (Dropout)	(None, 50)	0
dense (Dense)	(None, 25)	1,275
dense_1 (Dense)	(None, 1)	26

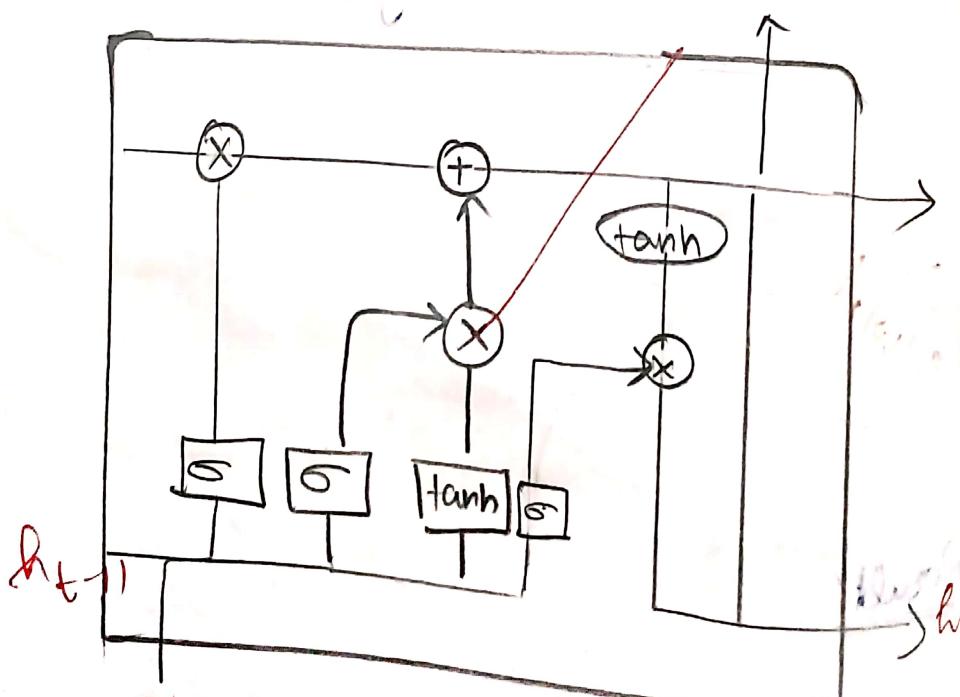
Total params: 31,901

Trainable params: 31,901

Non-Trainable params: 0

Training data shape: (946, 60, 1)

LSTM Architektur von hinten



5. Create Time series sequences.
for each epoch point in dataset
Input (x): Prev 60 closing prices.
Output (y): next closing price
6. Reshape Input data.
7. Build LSTM model.
8. Train the model.
9. Evaluate the model.
10. Visualise the results.
11. Analyse performance
12. STOP.

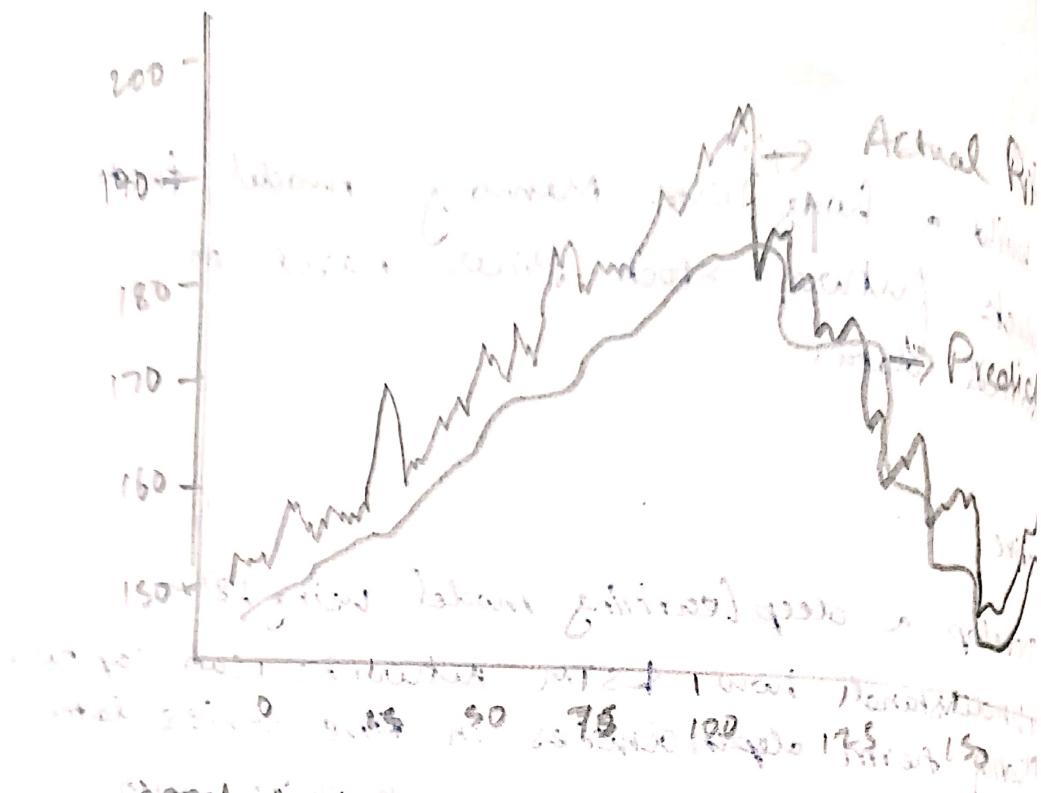
Observation:

1. Google's stock prices for past 5 years were successfully loaded & visualised
2. Preprocessing such as normalisation (minmaxscaler)
3. During training both RNN & LSTM demonstrated decreasing training loss
4. LSTM model showed smoother & more stable convergence.
 MSE was lower than RNN.

Result:

The code was successfully executed & verified

Model 1 - 100 Epochs



epoch 1 : loss : 0.8489

epoch 2 : loss : 0.000416

epoch 3 : loss : 0.0025

epoch 4 : loss : 0.0025

epoch 5 : loss : 0.0026

epoch 6 : loss : 0.0027

epoch 7 : loss : 0.0019

epoch 8 : loss : 0.0018

epoch 9 : loss : 0.0017

epoch 10 : loss : 0.0015

epoch 11 : loss : 0.0016

epoch 12 : loss : 0.0016

epoch 13 : loss : 0.0016

epoch 14 : loss : 0.0013

epoch 15 : loss : 0.0013

epoch 16 : loss : 0.0013

epoch 17 : loss : 0.0012

epoch 18 : loss : 0.0014

Q. Implementation Of Recurrent Neural Network

Aim

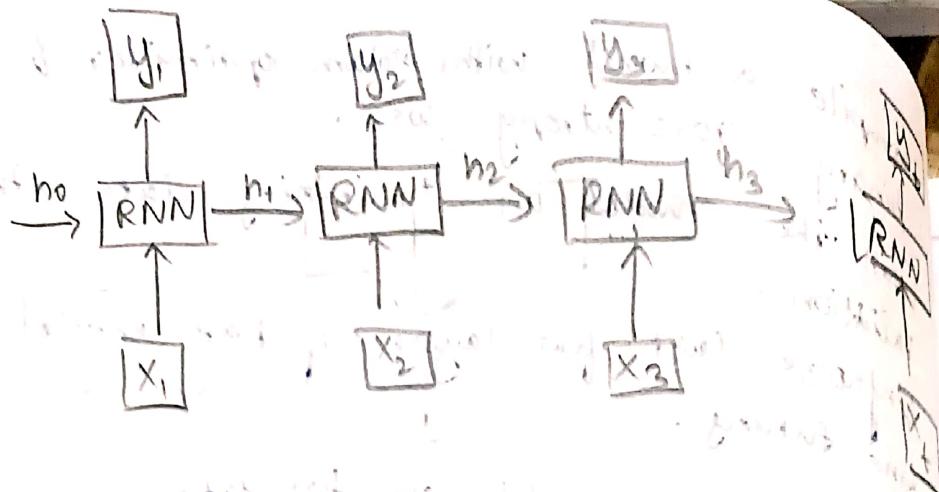
To design and implement a Recurrent Neural Network model for text classification and evaluate its performance using accuracy & confusion matrix.

Objectives:

1. To process ~~textual~~^{Stock} data into numerical sequences.
2. To build RNN model for analysing sequential data.
3. To train & validate the model for sentiment prediction.
4. To evaluate the model using accuracy, precision, recall, F1-score & confusion matrix.

Algorithm:

1. Import Necessary libraries
~~Tensorflow, Keras, Numpy etc..~~
2. Load of Preprocess the dataset.
~~Tokenisation, padding, train + test split~~
3. Build the RNN model using Keras Sequential API
4. Compile the model with an optimiser and loss function
5. Train the model on the training dataset
6. Evaluate the model using test data.



a. Implementation of RNN

Aim

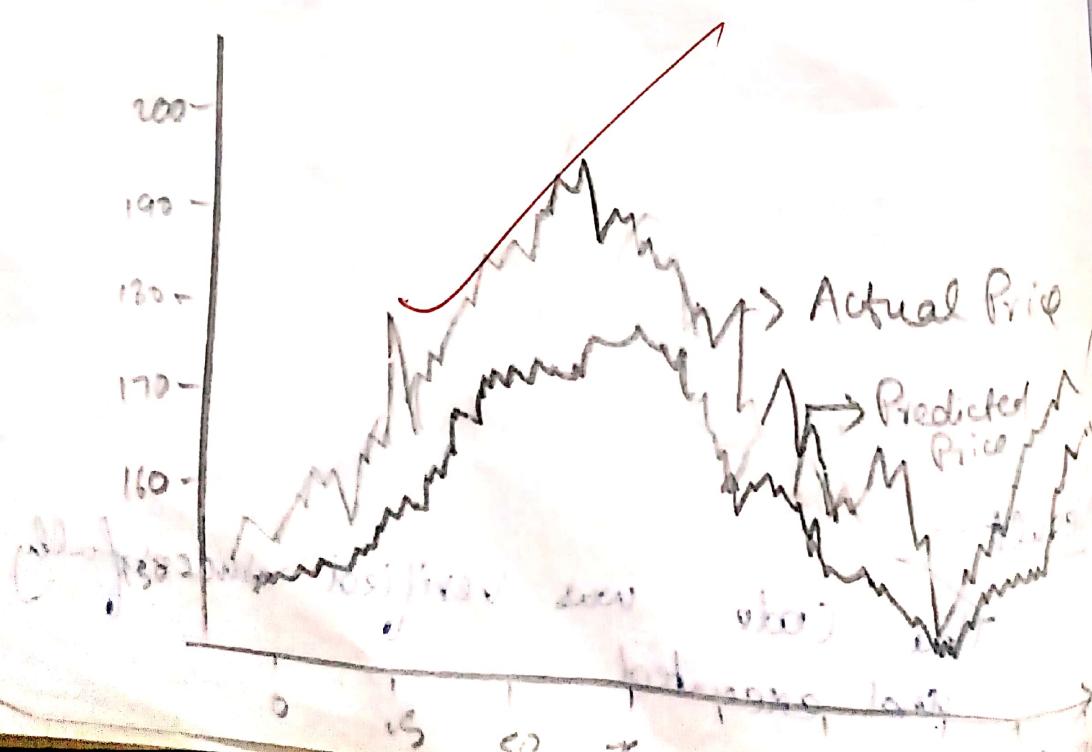
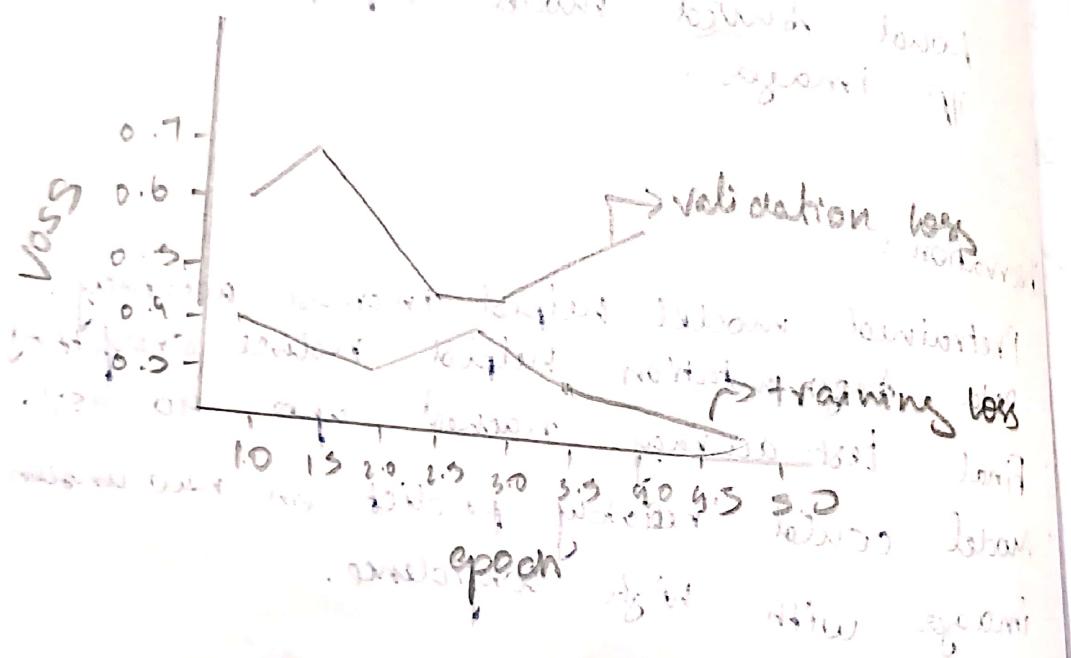
To design and implement a Neural Network model and evaluate its performance using confusion matrix.

Objectives:

1. To process Stop words.
2. To build RNN data.
3. To train + validate the sentiment prediction model.
4. To evaluate precision, recall and F1 score.

Algorithm:

1. Import Necessary Libraries
2. Load & Preprocess Data
3. Build the Model
4. Compile Model and Loss Function
5. Train the Model
6. Evaluate the Model



Epoch 1 : 1 minute

loss : 0.5258 accuracy : 0.7373

Val. loss : 0.4582 - val accuracy : 0.80

Epoch 2 : 1 minute

loss : 0.3909 acc : 0.8317 val loss : 0.4121 Val

: Epoch 3 : 1 minute

loss : 0.3168 acc : 0.8692 val loss : 0.3896 Val

Epoch 4

loss : 0.2657 acc : 0.8923 val loss : 0.4054 Val

Epoch 5 : 1 minute

loss : 0.2263 acc : 0.9102 val loss : 0.4458 Val

Test loss : 0.4127

Test Acc : 0.8174

Precision recall f1-score supp

Precision	recall	f1-score	supp
0.70	0.69	0.68	0.23
0.73	0.75	0.74	0.23

macro avg

precision	recall	f1-score	supp
0.717	0.717	0.717	500

Mean Squared Error : 1.5976930735

Mean Absolute error : 6.5673052182



7. Generate the confusion matrix and Classification report.

8. Display accuracy & loss graphs.

Observation:

If RNN has short term memory, so it forget informations to retain its memories.

Sometimes in sequential data the continuous values like loss in stock prices is crucial to predict future prices.

This is vanishing gradient problem.

This is resolved by LSTM.

~~BiLSTM~~

Result:

The code was successful written & executed.

Perform Compression on MNIST DATASET Using Auto Encoder.

Aim:

To implement an autoencoder using pytorch for compressing and reconstructing images from the MNIST dataset, demonstrating unsupervised feature learning and dimensionality reduction.

Objectives:

1. To understand the working principle of an autoencoder and its encoder - decoder structure.
2. To build and train a fully connected autoencoder using Pytorch on the mnist dataset.
3. To evaluate the models ability to compress and reconstruct images.
4. To visualise the original and reconstructed image to analyse reconstruction quality

Pseudocode

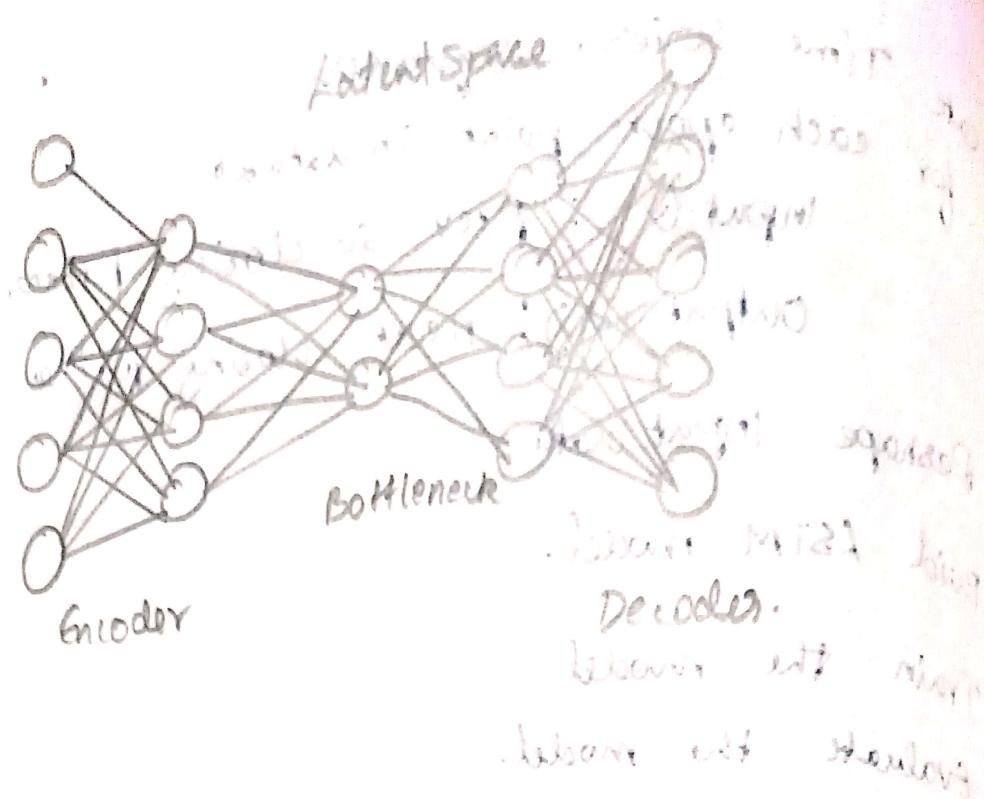
BEGIN

 Import torch, torch.nn, torch.optim
 Load MNIST dataset

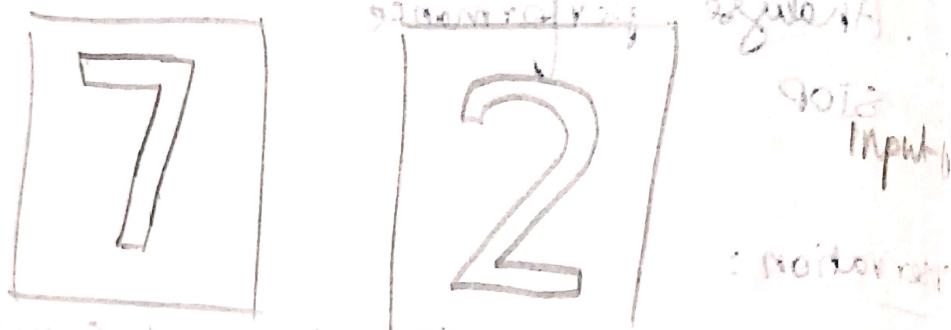
 Normalise images to $[0, 1]$

 Define encoder:

 Input \rightarrow Dense (128, relu) \rightarrow Dense (64, relu)
 Dense (32, relu)



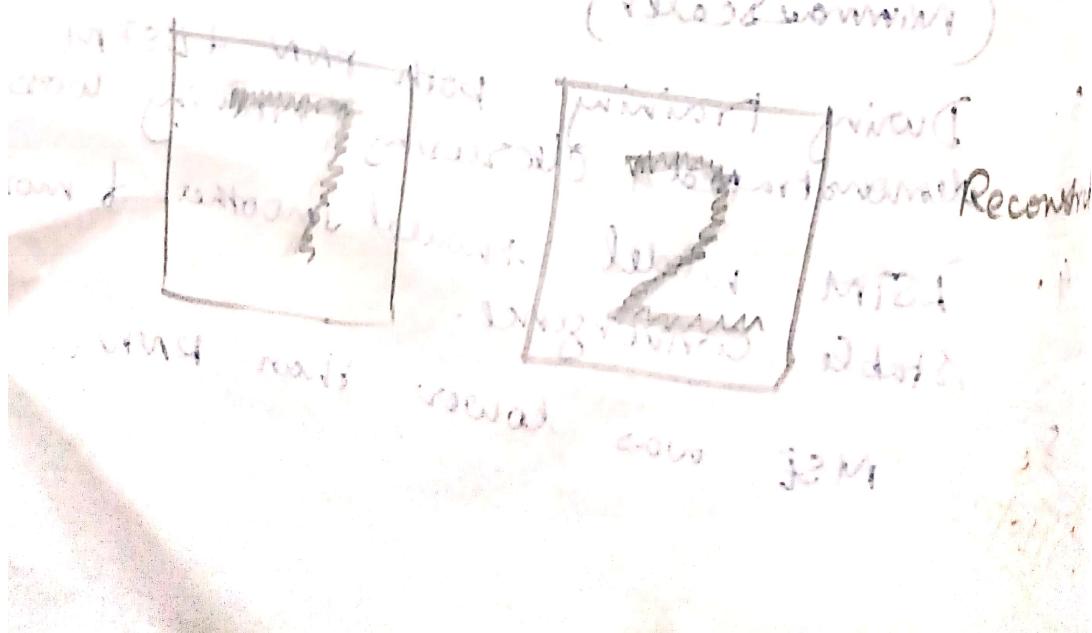
Original Image: Number 2



Decoding rule of latent space & input
behavior & behavior of neurons now

Reconstructed Image.

(This is working)



Define decoder

Dense (64, relu) \rightarrow Dense(128, relu) \rightarrow Dense(784, sigmoid)

Combine encoder and decoder \rightarrow Autoencoder.

Compile (Autoencoder, optimizer = 'adam', loss=mae)

Train on MNIST Images (input = output)

Visualise original vs reconstructed Images

Observation

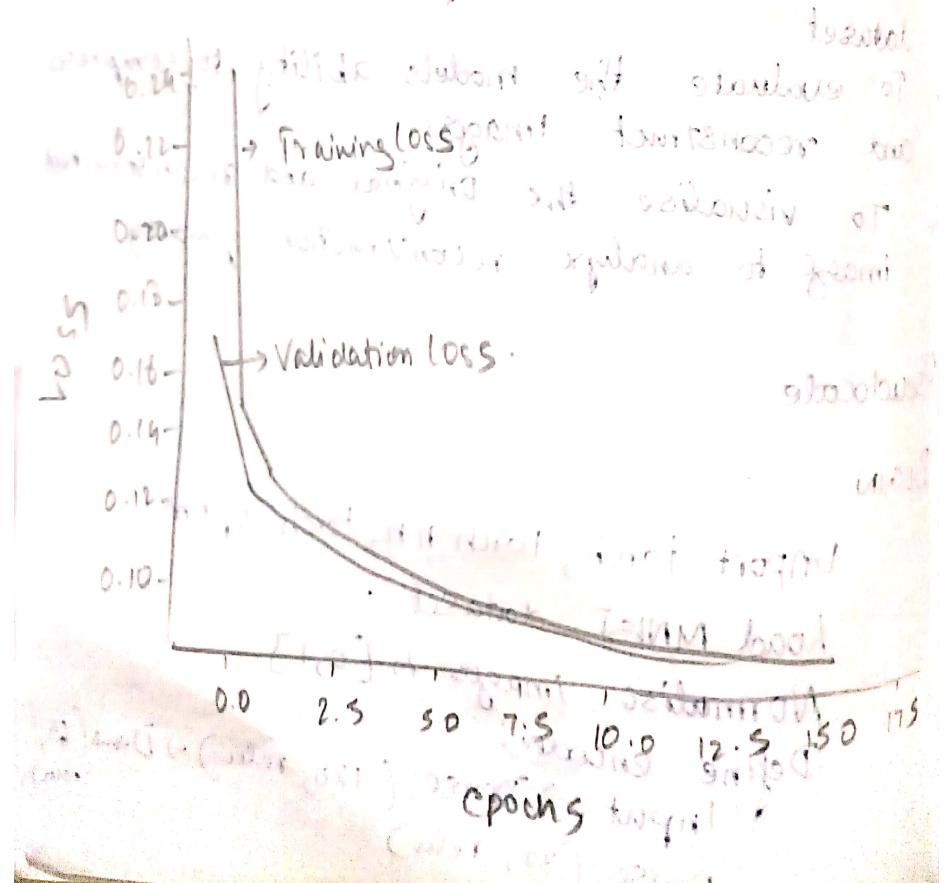
As epochs increases, reconstruction loss decreases, and the model learns compressed latent feature (32-dim representation).

Ans

Result:

The Autoencoder successfully compressed and reconstructed MNIST Images with minimal loss.

Training data shape (60000, 784)
 Test data shape (10000, 784)
 Epoch Val-loss
 0 105.5 16.67
 1 0.3416 6.1329
 2 0.1950 0.1222
 3 0.1309 0.1533
 4 0.1213 0.1103
 5 0.1149 0.0973
 6 0.1107 0.0973
 7 0.1025 0.1043
 8 0.1025 0.1043
 9 0.1033 0.1004
 10 0.1012 0.0980
 11 0.0999 0.0975
 12 0.0986 0.0967
 13 0.0974 0.0957
 14 0.0963 0.0963
 15 0.0957 0.0941
 16 0.0941 0.0941



Experiment Using Variational Autoencoder (VAE)

Aim:

To implement a Variational Autoencoder using pytorch for learning probabilities latent representation and generating new handwritten digital images from MNIST dataset.

Objective

1. To understand the concept of VAE and how they differ from standard autoencoder.
2. To implement VAE model using pytorch that learn a latent distribution.
3. To reconstruct and generate new images using the learned latent space.
4. To evaluate the generative capability of the trained VAE.

Pseudocode:

BEGIN

Import torch, torch nn, torch optim

load MNIST data.

Define VAE C

Encoder

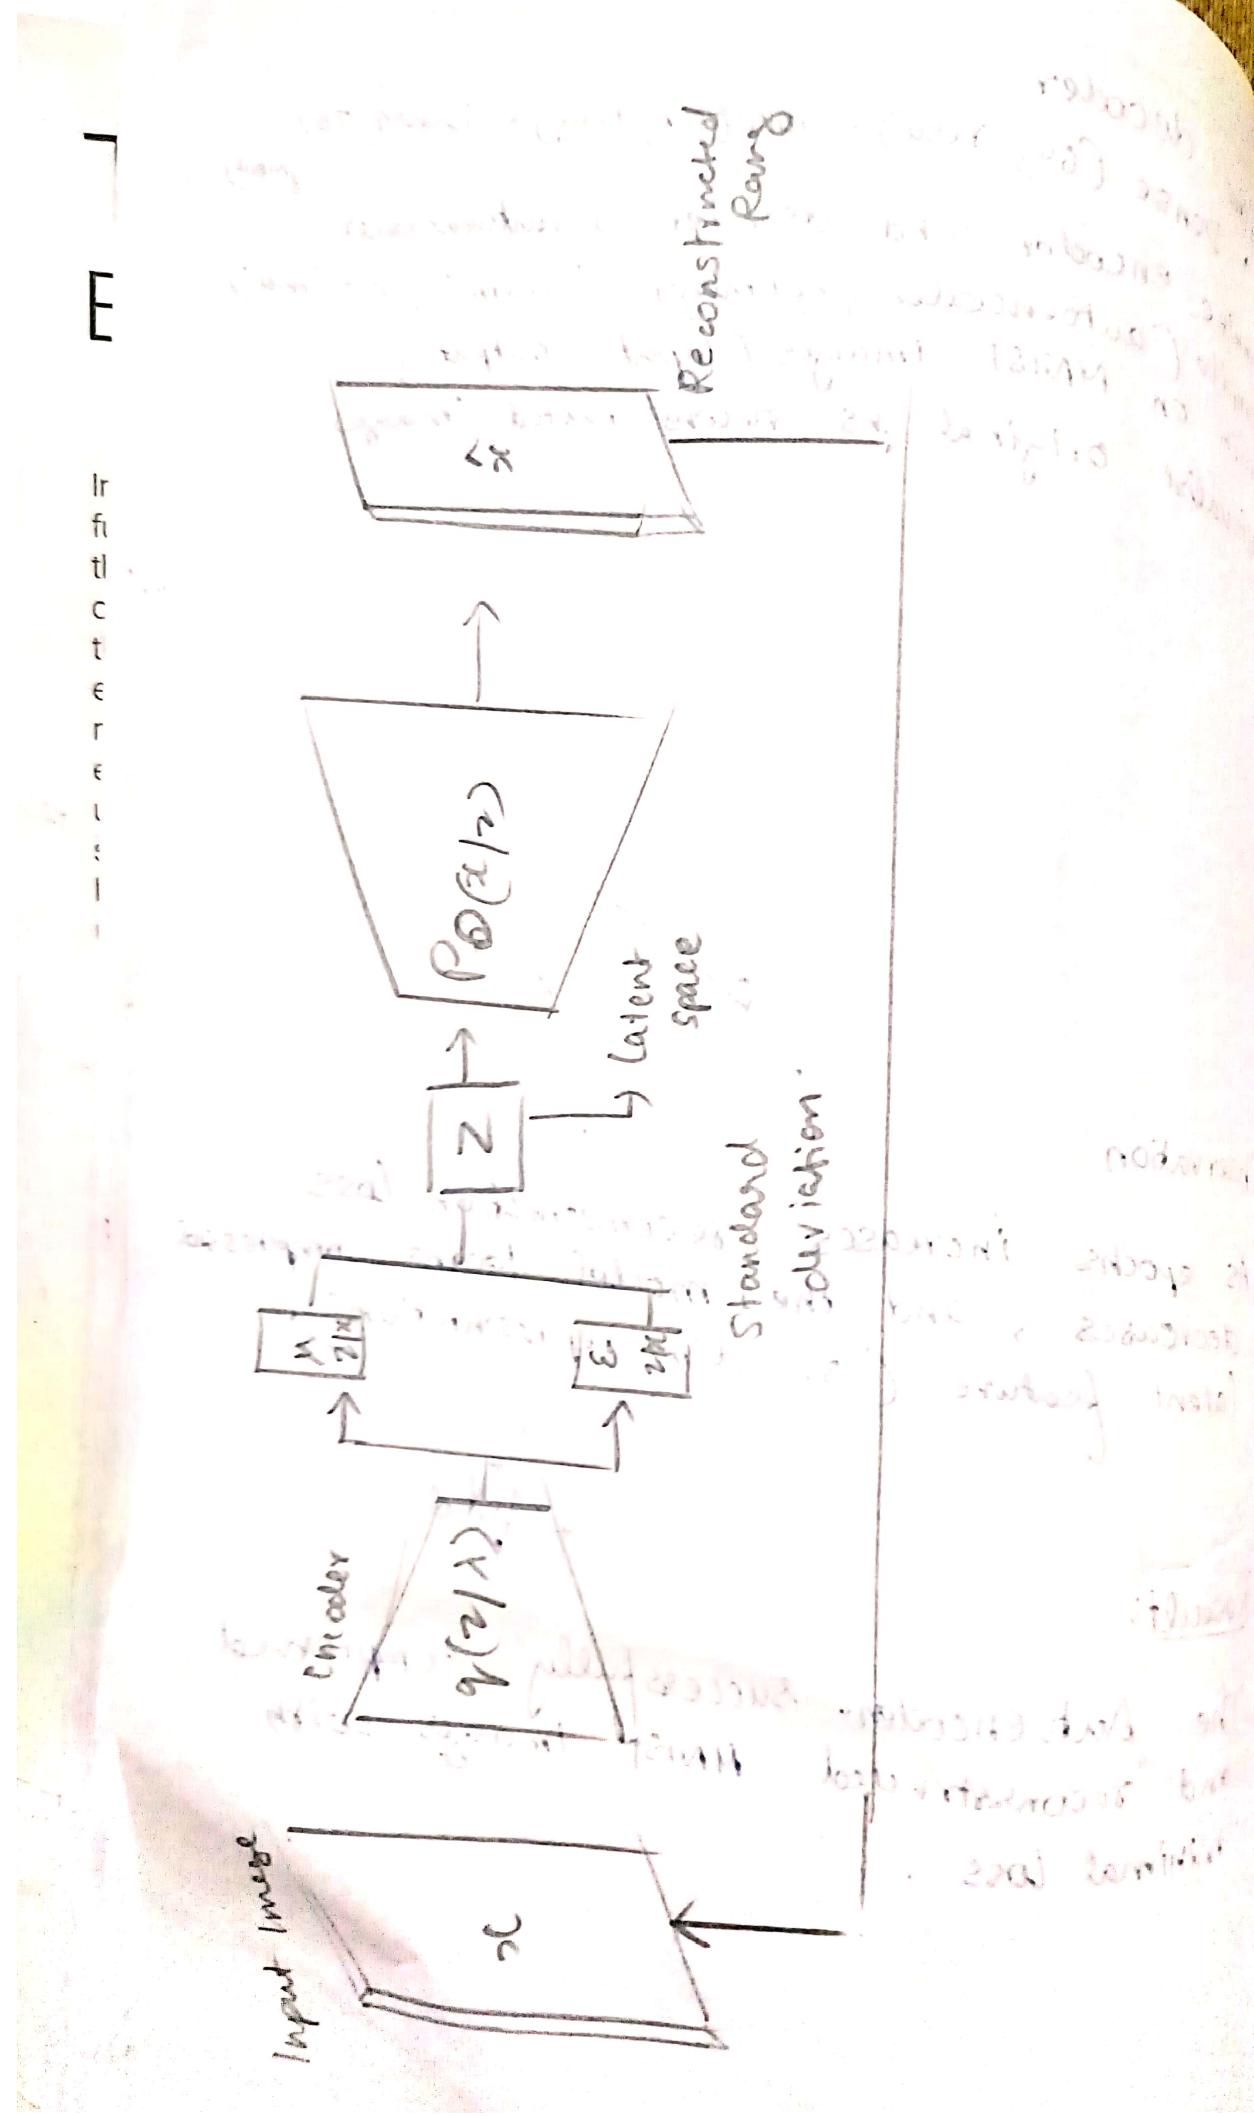
Reparameterisation

Decoder

Forward

Pefine loss function

Reconstruct loss.



Initialise model, optimiser(Adam)

For each epoch :

 For epoch each batch in training data :

 For recon, $\text{KL} = \text{mod}(f(\text{input}))$

 loss = BCE + KL

 Back propagation and update weights

 Print Average epoch losses.

Test model

Display generated samples.

END.

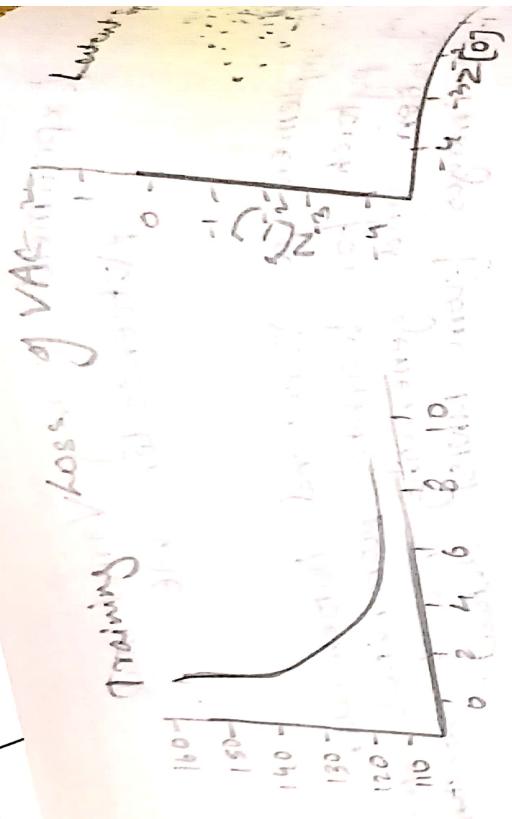
Observation:

Generated Images resemble digits but with some variations due to Sampling from latent space.



Result:

The VAE generated new digit images by learning probabilistic latent distributions, validating generative capability.



Epoch	KL-loss	recons+recThr loss	Val KL-loss
1	10.07931	257.8823	4.0952
2	4.91999	770.6052	4.9311
3	4.82984	161.5327	5.0213
4	5.11124	157.6767	5.2313
5	5.3440	156.0489	5.4457
6	5.5002	151.4706	5.3396
7	5.5985	149.8563	5.5174
8	5.6778	148.5911	5.6217
9	5.7714	147.2466	5.8601

30 1.4627 136.0784 1.7 6.3085

13 Implement Deep Convolutional GAN to generate Complex Color images.

Aim:

To implement a Deep Convolutional Generative Adversarial Network for generating Complex color images.

Objective

- To understand adversarial learning between Generator and discriminator.
- To generate realistic color images.
- To observe the improvement of generated images over training epochs.
- To study latent space representation and image synthesis.

PseudoCode:

Load and normalise color image dataset

Define Generator:

Dense \rightarrow Reshape \rightarrow Conv2D Transpose layers
(ReLU, Batch Norm) \rightarrow Output(3-channels, tanh)

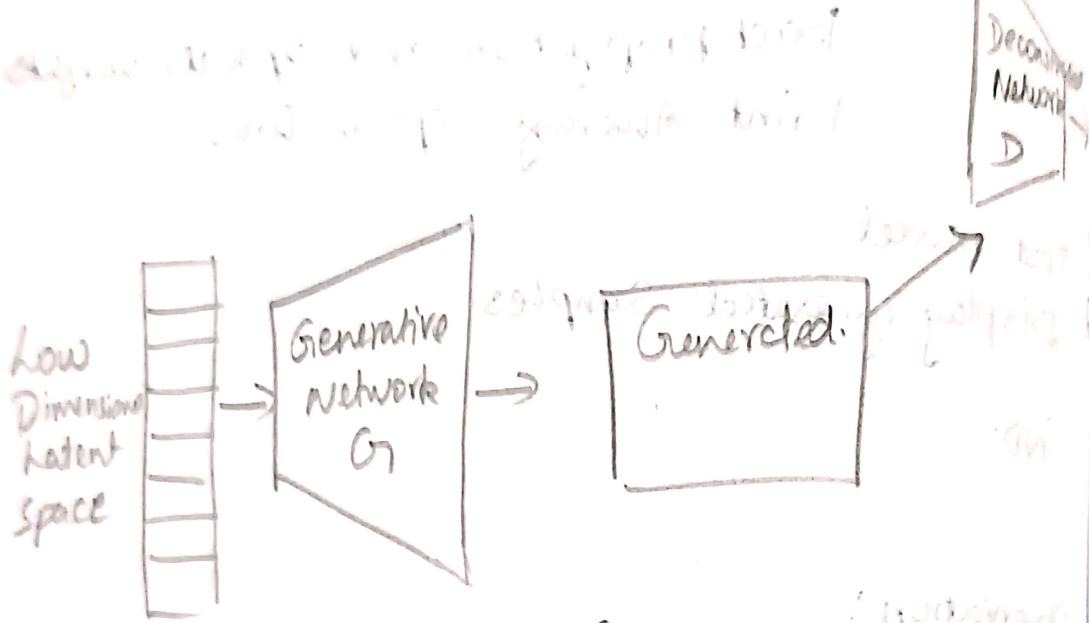
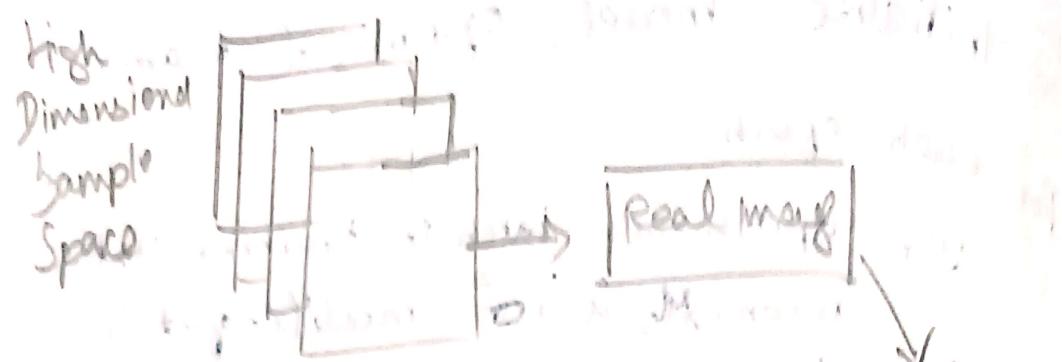
Define Discriminator:

Conv2D layers (Leaky ReLU, Dropout) \rightarrow
Dense(1, sigmoid)

Train adversarially

for each epoch:

Train Discriminator on real + fake images

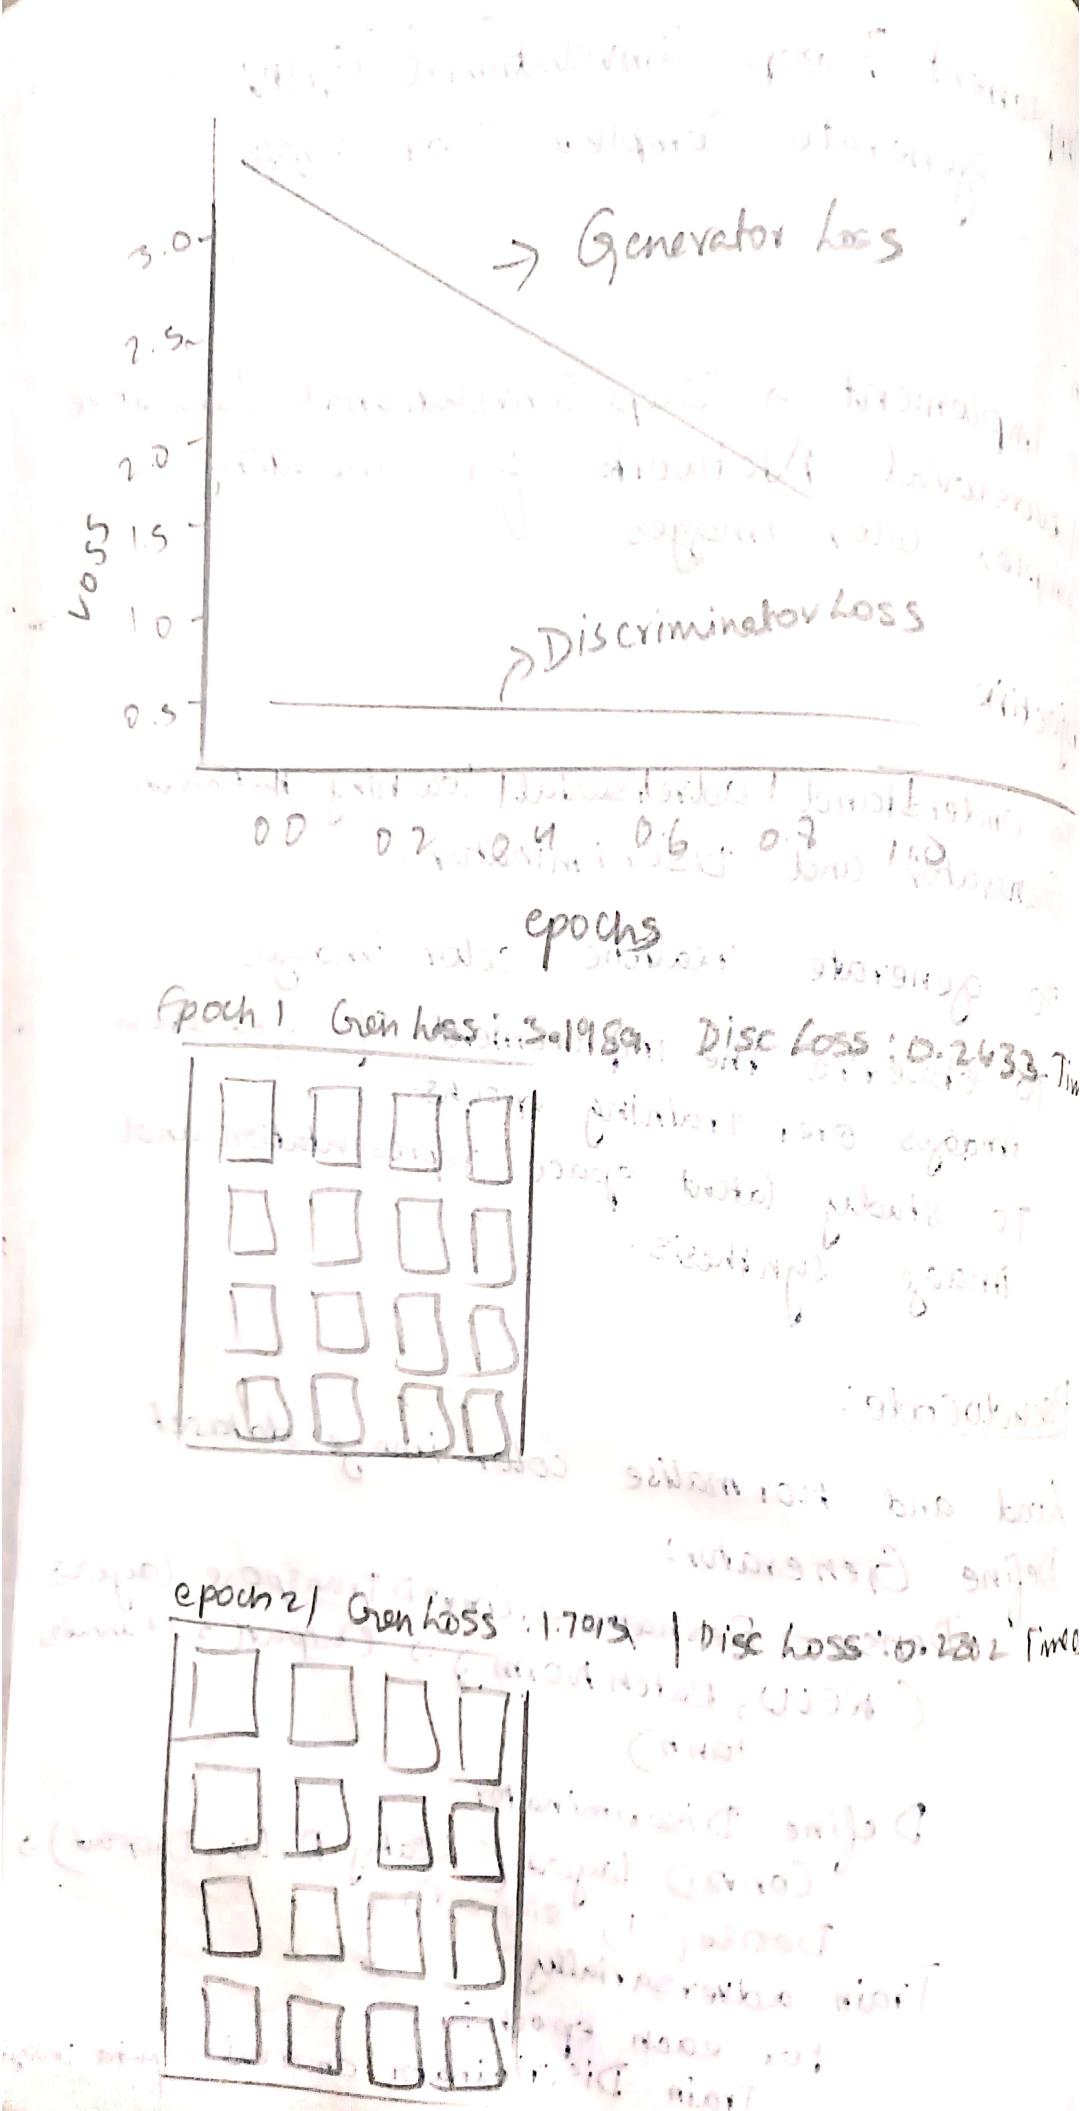


GAN Architecture

This is how Generative Adversarial Networks function at the architecture level.



The latent space is used to generate training data for the discriminative network. This training data is then used to train the generative network.



Train Generator via combined model
(Generator + frozen Discriminator)
Generate and visualise fake images.

Observation

Initially, images are noisy, after several epochs, generator starts producing realistic structures and colors.

Result

The Code was executed & Verified successfully.

Understanding the architecture of Pre-trained Model.

Aim:

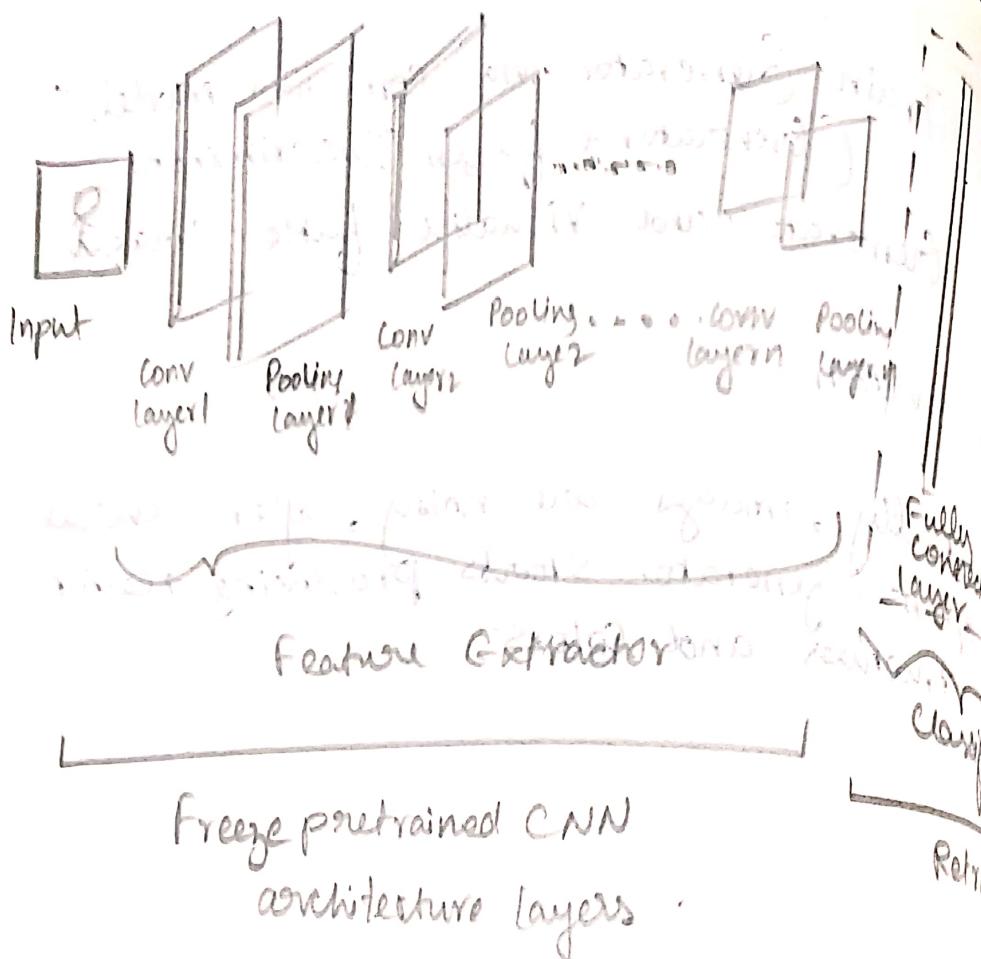
To study and analyse the architecture of pre-trained CNN models like ResNet and Inception.

Objective:

- To understand the layer-wise design of state-of-the-art model.
- To visualise feature extraction process in deep CNNs.
- To analyse parameter distribution and feature map visualisation.
- To explore how pre-trained models are used for transfer learning tasks.

Pseudo code

1. Import necessary libraries (Tensorflow/Keras or pytorch)
2. Load a Pre-trained CNN model.
with ImageNet weights.
3. Exclude top layers if needed.
4. Display and study the model summary to understand layers and parameters.
5. Load & preprocess a sample image.
6. Pass the image through the model to visualise feature maps.



	precision	recall	f1-score	Support
airplane	0.33	0.66	0.46	103
automobile	0.61	0.36	0.38	89
bird	0.32	0.12	0.18	100
cat	0.60	0.03	0.06	103
deer	0.24	0.34	0.28	90
dog	0.30	0.03	0.06	86
frog	0.26	0.70	0.38	112
horse	0.43	0.25	0.33	102
ship	0.53	0.22	0.31	106
truck	0.42	0.54	0.48	109

Accuracy

Macro avg 0.34 0.34 0.34 100%

Weighted avg 0.34 0.34 0.34 100%

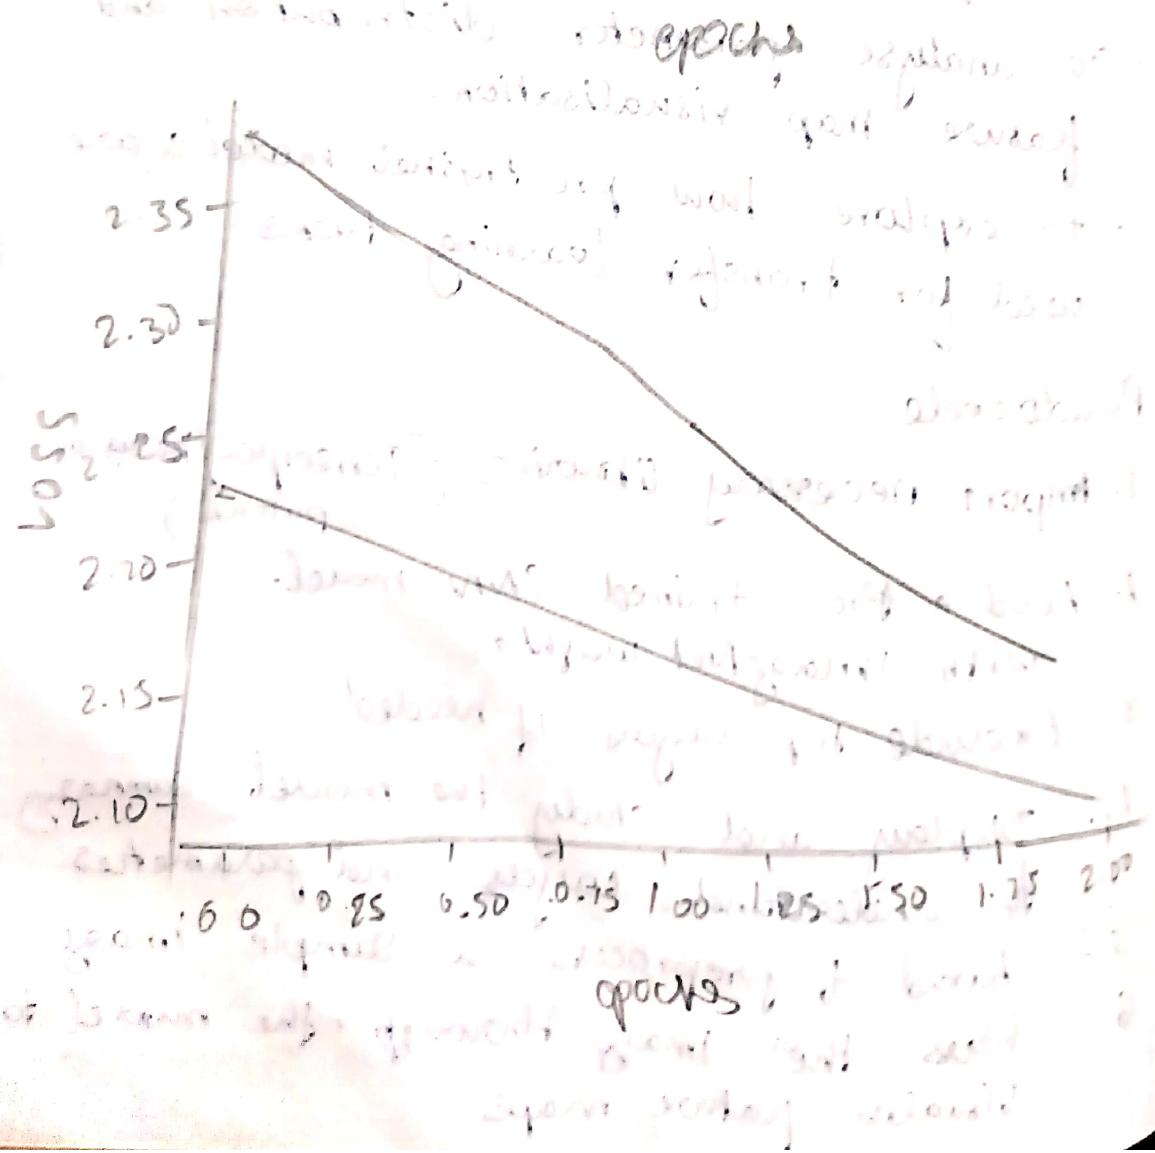
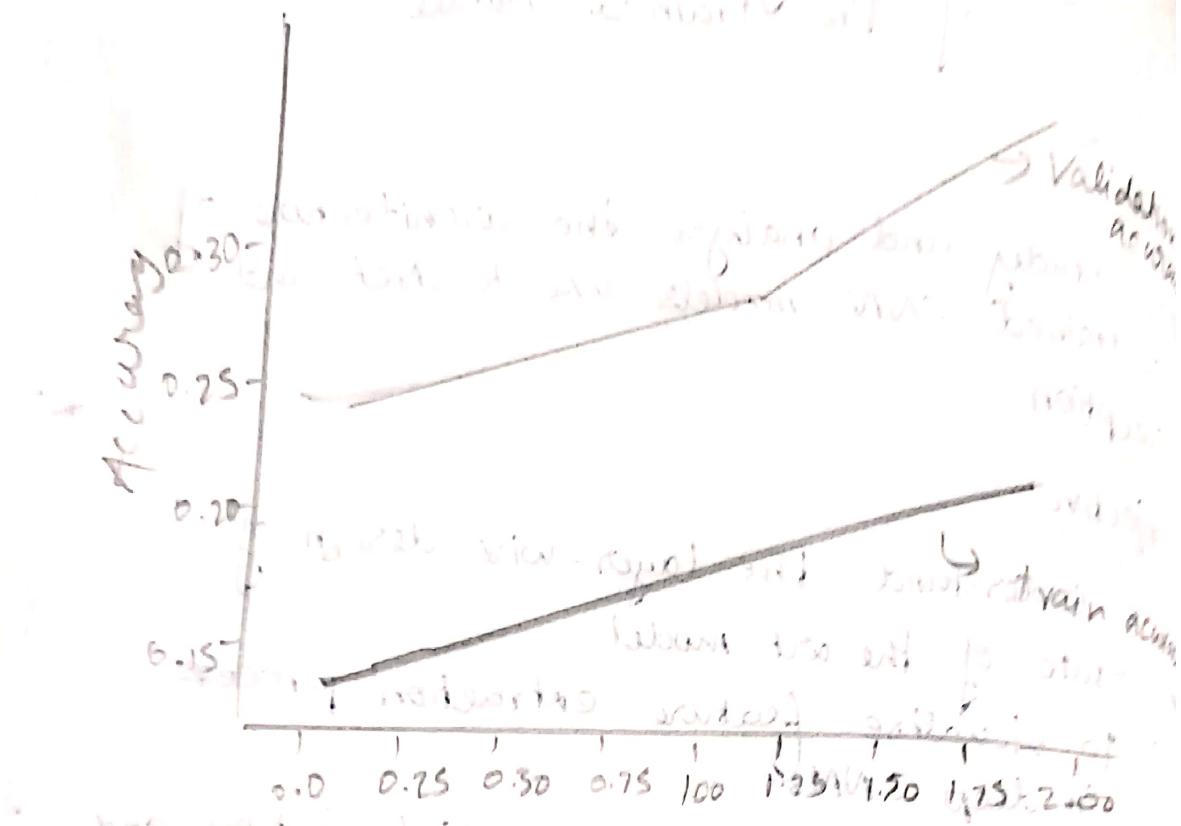
7. Analyse how different layers extract various levels of features

Observation

1. The model summary displayed multiple convolutional, pooling, and fully connected layers arranged in a hierarchical structure.
2. The early layers captured low-level features such as edges, textures, and simple patterns.
3. The deeper layers extracted high level abstract features like object parts and shapes.
4. Feature map visualisation showed that each filter responds to specific image characteristics.
5. Models like VGG16 had large no. of parameters with uniform layer depth, while ResNet used skip connection to improve learning.
6. The pre-trained model demonstrated efficient feature extraction, confirming its ability to generalise to various image types.

Result:

The code was successfully executed and verified.



Implement a Pre trained CNN Model as a feature extractor Using Transfer learning -

Aim:

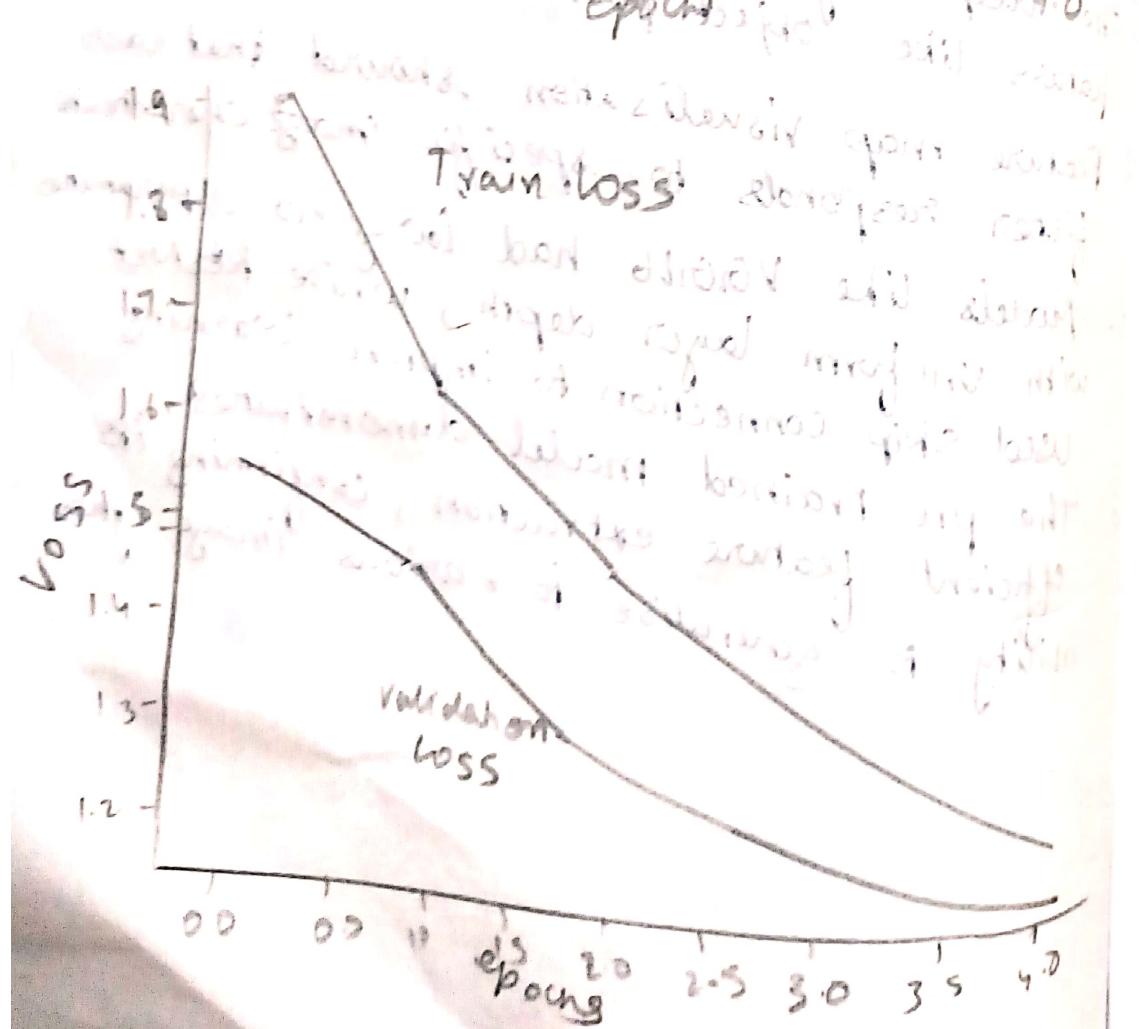
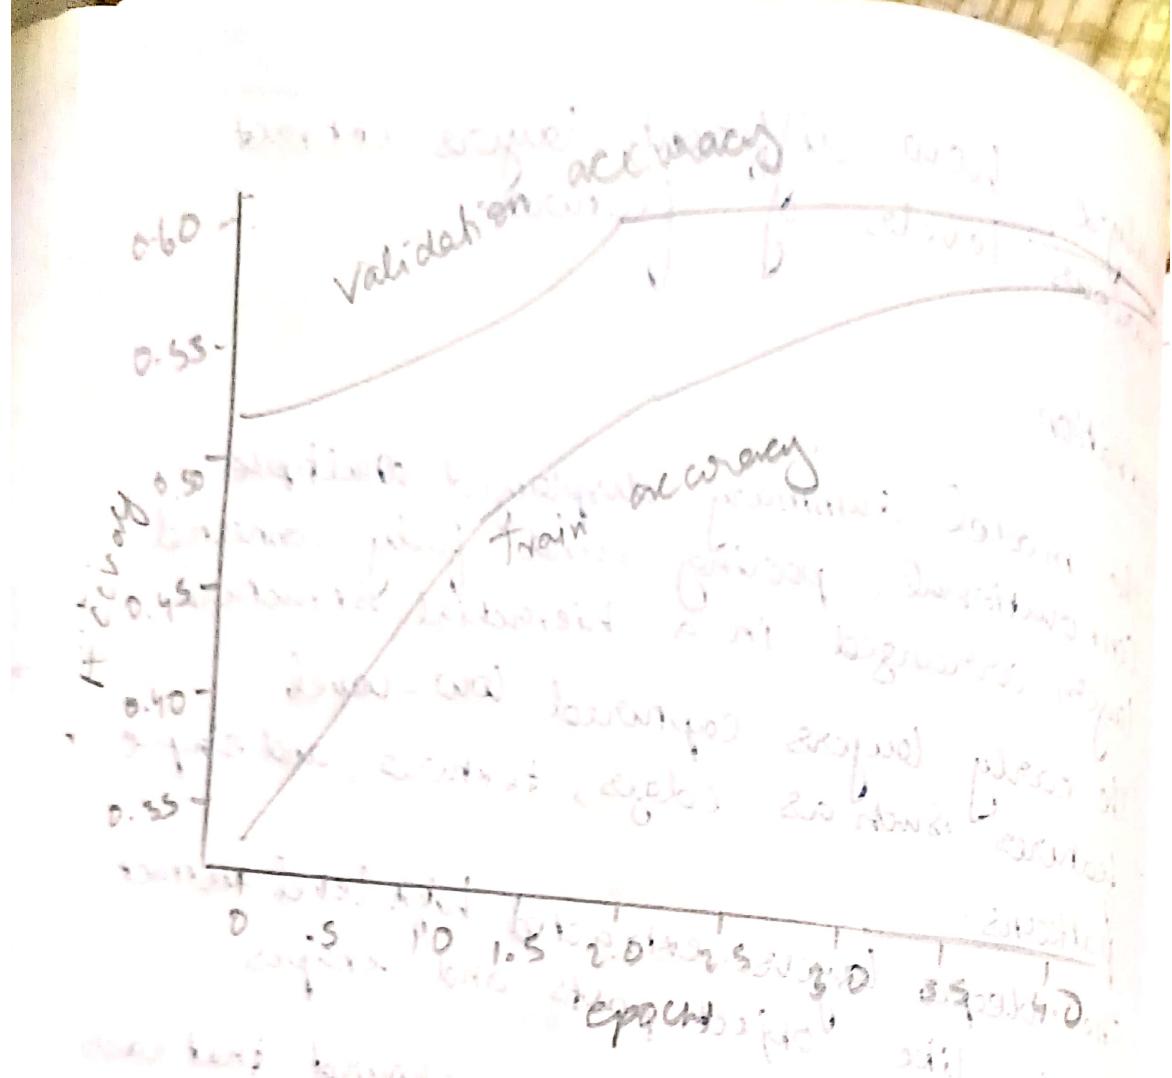
To Implement a pre-trained CNN model as a feature extractor using transfer learning for image classification .

Objectives:

1. To reuse pretrained CNN weights for a new dataset.
2. To extract high level image features from convolutional layers.
3. To fine-tune the model for better performance on custom tasks.
4. To reduce training time and computational cost through transfer learning .

Pseudo Code:

1. Import required libraries and load a pre-trained CNN model
2. Freeze the convolutional base layers to preserve learned features.
3. Add custom layers (flatten -> dense -> output layers with softmax)
4. Compile the model with an optimizer like Adam and suitable loss function.
5. Load and preprocess the custom dataset.
6. Train the model on new data for limited epochs.



epoch	accuracy	loss	Val'accuracy
1	0.2383	2.613	0.5150
2	0.4471	1.5875	0.5440
3	0.5238	1.3993	0.6010
4	0.5606	1.2921	0.6060
5	0.6146	1.1975	0.6060

Total execute time : 96.27 seconds.

* obj detection uses bounding box of
part of car and zoom level digital zooming up
and down to detect objects even
if they are small.
* detection of license plate is done by
using neural network trained on
large amount of images.
* In case of license plate detection
it uses bounding box of image
and crop part of image which contains
license plate and then
detects license plate from it.



7. Evaluate the model Accuracy and visualise training performance.

Observation:

1. The model converged quickly with fewer epochs due to the pre-trained feature extractor.
2. The lower convolutional layers captured general features, while the new dense layers learned task specific patterns.
3. Validation accuracy improved steadily without overfitting.
4. Training and loss curves showed stable learning behaviour.

Result:

The code was successfully verified & executed.

Implement a YOLO Model to Detect Objects.

Aim:

To implement a YOLO model for real time Object detection in images / video streams.

Objectives:

1. To Understand the working of one stage Object detection using YOLO architecture.
2. To detect and classify multiple objects within an image simultaneously.
3. To visualise bounding boxes & confidence Scores for detected objects.
4. To evaluate the model's speed & accuracy in real time detection tasks.

Pseudocode:

1. Import required libraries (openCV, Tensorflow, Pytorch)
2. Load the pre-trained YOLO model and class labels.
3. Load the input image or video streams.
4. Preprocess the frame (resize, normalize, convert to tensor)
5. Perform detection using YOLO model.
6. Extract bounding box coordinates, class names, and confidence scores.

7. Draw bounding boxes and labels on the image.
8. Display or save the output image/video with detected objects.

Observation:

1. The model successfully detected multiple object with bounding boxes and labels.
2. Detection speed was fast, suitable for real time application.
3. Confidence scores varied based on object clarity and lighting.
4. Bounding boxes accurately localized the detected objects in most cases.

Result:

The code was successfully executed and verified.